

## chap4\_exercise\_labs

```
library(ISLR2)
library(ggplot2)
library(MASS)
```

Attaching package: 'MASS'

The following object is masked from 'package:ISLR2':

Boston

```
library(e1071)
library(class)
library(dplyr)
```

Attaching package: 'dplyr'

The following object is masked from 'package:MASS':

select

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

#####labs

head(Smarket)

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
1	2001	0.381	-0.192	-2.624	-1.055	5.010	1.1913	0.959	Up
2	2001	0.959	0.381	-0.192	-2.624	-1.055	1.2965	1.032	Up
3	2001	1.032	0.959	0.381	-0.192	-2.624	1.4112	-0.623	Down
4	2001	-0.623	1.032	0.959	0.381	-0.192	1.2760	0.614	Up
5	2001	0.614	-0.623	1.032	0.959	0.381	1.2057	0.213	Up
6	2001	0.213	0.614	-0.623	1.032	0.959	1.3491	1.392	Up

names(Smarket)

[1]	"Year"	"Lag1"	"Lag2"	"Lag3"	"Lag4"	"Lag5"
[7]	"Volume"	"Today"	"Direction"			

summary(Smarket)

Year		Lag1		Lag2		Lag3	
Min.	:2001	Min.	:-4.922000	Min.	:-4.922000	Min.	:-4.922000
1st Qu.:	:2002	1st Qu.:	-0.639500	1st Qu.:	-0.639500	1st Qu.:	-0.640000
Median	:2003	Median	: 0.039000	Median	: 0.039000	Median	: 0.038500
Mean	:2003	Mean	: 0.003834	Mean	: 0.003919	Mean	: 0.001716
3rd Qu.:	:2004	3rd Qu.:	0.596750	3rd Qu.:	0.596750	3rd Qu.:	0.596750
Max.	:2005	Max.	: 5.733000	Max.	: 5.733000	Max.	: 5.733000
Lag4		Lag5		Volume		Today	
Min.	:-4.922000	Min.	:-4.92200	Min.	:0.3561	Min.	:-4.922000
1st Qu.:	:-0.640000	1st Qu.:	-0.64000	1st Qu.:	:1.2574	1st Qu.:	-0.639500
Median	: 0.038500	Median	: 0.03850	Median	:1.4229	Median	: 0.038500
Mean	: 0.001636	Mean	: 0.00561	Mean	:1.4783	Mean	: 0.003138
3rd Qu.:	0.596750	3rd Qu.:	0.59700	3rd Qu.:	:1.6417	3rd Qu.:	0.596750
Max.	: 5.733000	Max.	: 5.73300	Max.	:3.1525	Max.	: 5.733000
Direction							
Down:602							

Up :648


```
cor(Smarket[, -9])
```

	Year	Lag1	Lag2	Lag3	Lag4
Year	1.00000000	0.029699649	0.030596422	0.033194581	0.035688718
Lag1	0.02969965	1.000000000	-0.026294328	-0.010803402	-0.002985911
Lag2	0.03059642	-0.026294328	1.000000000	-0.025896670	-0.010853533
Lag3	0.03319458	-0.010803402	-0.025896670	1.000000000	-0.024051036
Lag4	0.03568872	-0.002985911	-0.010853533	-0.024051036	1.000000000
Lag5	0.02978799	-0.005674606	-0.003557949	-0.018808338	-0.027083641
Volume	0.53900647	0.040909908	-0.043383215	-0.041823686	-0.048414246
Today	0.03009523	-0.026155045	-0.010250033	-0.002447647	-0.006899527

	Lag5	Volume	Today
Year	0.029787995	0.53900647	0.030095229
Lag1	-0.005674606	0.04090991	-0.026155045
Lag2	-0.003557949	-0.04338321	-0.010250033
Lag3	-0.018808338	-0.04182369	-0.002447647
Lag4	-0.027083641	-0.04841425	-0.006899527
Lag5	1.000000000	-0.02200231	-0.034860083
Volume	-0.022002315	1.00000000	0.014591823
Today	-0.034860083	0.01459182	1.000000000

```
attach(Smarket)
plot(Volume)
```



#Logistic Regression

```
glm.fits <- glm(
  Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,
  data = Smarket,
  family = binomial
)
summary(glm.fits)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
     Volume, family = binomial, data = Smarket)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-0.126000	0.240736	-0.523	0.601
Lag1	-0.073074	0.050167	-1.457	0.145
Lag2	-0.042301	0.050086	-0.845	0.398
Lag3	0.011085	0.049939	0.222	0.824
Lag4	0.009359	0.049974	0.187	0.851
Lag5	0.010313	0.049511	0.208	0.835

```
Volume      0.135441  0.158360  0.855  0.392
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 1731.2  on 1249  degrees of freedom
Residual deviance: 1727.6  on 1243  degrees of freedom
AIC: 1741.6
```

Number of Fisher Scoring iterations: 3

**smallest p-value here is associated with Lag1**

```
coef(glm.fits)
```

```
(Intercept)      Lag1      Lag2      Lag3      Lag4      Lag5
-0.126000257 -0.073073746 -0.042301344  0.011085108  0.009358938  0.010313068
      Volume
0.135440659
```

```
summary(glm.fits)$coef
```

```
              Estimate Std. Error    z value Pr(>|z|)
(Intercept) -0.126000257 0.24073574 -0.5233966 0.6006983
Lag1         -0.073073746 0.05016739 -1.4565986 0.1452272
Lag2         -0.042301344 0.05008605 -0.8445733 0.3983491
Lag3          0.011085108 0.04993854  0.2219750 0.8243333
Lag4          0.009358938 0.04997413  0.1872757 0.8514445
Lag5          0.010313068 0.04951146  0.2082966 0.8349974
Volume       0.135440659 0.15835970  0.8552723 0.3924004
```

```
summary(glm.fits)$coef[, 4]###col 5 the p-value
```

```
(Intercept)      Lag1      Lag2      Lag3      Lag4      Lag5
0.6006983  0.1452272  0.3983491  0.8243333  0.8514445  0.8349974
      Volume
0.3924004
```

```
glm.probs <- predict(glm.fits, type = "response")
glm.probs[1:10]
```

```

      1      2      3      4      5      6      7      8
0.5070841 0.4814679 0.4811388 0.5152224 0.5107812 0.5069565 0.4926509 0.5092292
      9     10
0.5176135 0.4888378
```

```
contrasts(Direction)
```

```

      Up
Down  0
Up    1
```

```
glm.pred <- rep("Down", 1250)
glm.pred[glm.probs > .5] = "Up"
```

```
table(glm.pred, Direction)
```

```

      Direction
glm.pred Down  Up
Down    145 141
Up      457 507
```

```
mean(glm.pred == Direction)
```

```
[1] 0.5216
```

```
train <- (Year < 2005)
Smarket.2005 <- Smarket[!train, ]#test df
dim(Smarket.2005)
```

```
[1] 252  9
```

```
Direction.2005 <- Direction[!train]
```

```
glm.fits <- glm(  
  Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,  
  data = Smarket, family = binomial, subset = train  
)
```

#use the test\_x to predict

```
glm.probs <- predict(glm.fits, Smarket.2005,  
  type = "response")
```

```
glm.pred <- rep("Down", 252)  
glm.pred[glm.probs > .5] <- "Up"  
table(glm.pred, Direction.2005)#Direction.2005 truth/test y
```

```
      Direction.2005  
glm.pred Down Up  
Down    77 97  
Up      34 44
```

```
mean(glm.pred == Direction.2005)
```

```
[1] 0.4801587
```

#refit use Lag1 and Lag2

```
glm.fits <- glm(  
  Direction ~ Lag1 + Lag2 ,  
  data = Smarket, family = binomial, subset = train  
)
```

```
glm.probs <- predict(glm.fits, Smarket.2005,  
  type = "response")
```

```
glm.pred <- rep("Down", 252)  
glm.pred[glm.probs > .5] <- "Up"  
table(glm.pred, Direction.2005)
```

```

      Direction.2005
glm.pred Down  Up
Down    35   35
Up      76  106

```

```
mean(glm.pred == Direction.2005)
```

```
[1] 0.5595238
```

```

predict(glm.fits,
        newdata = data.frame(Lag1 = c(1.2, 1.5), Lag2 = c(1.1, -0.8)),
        type = "response" )

```

```

      1      2
0.4791462 0.4960939

```

```
#LDA
```

```

lda.fit <- lda(Direction ~ Lag1 + Lag2,
               data = Smarket,
               subset = train)

lda.fit

```

Call:

```
lda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
```

Prior probabilities of groups:

```

      Down      Up
0.491984 0.508016

```

Group means:

```

      Lag1      Lag2
Down 0.04279022 0.03389409
Up   -0.03954635 -0.03132544

```

Coefficients of linear discriminants:

```

      LD1
Lag1 -0.6420190
Lag2 -0.5135293

```



```
lda.pred <- predict(lda.fit, Smarket.2005)
lda.class <- lda.pred$class
table(lda.class, Direction.2005)#LDA,logistic almost identical
```

```

          Direction.2005
lda.class Down  Up
      Down   35  35
      Up    76 106
```

```
mean(lda.class == Direction.2005)
```

```
[1] 0.5595238
```

```
sum(lda.pred$posterior[, 1] >= .5)
```

```
[1] 70
```

```
sum(lda.pred$posterior[, 1] < .5)#the col2 is p(Down)
```

```
[1] 182
```

```
head(lda.pred$posterior)
```

```

          Down      Up
999  0.4901792 0.5098208
1000 0.4792185 0.5207815
1001 0.4668185 0.5331815
1002 0.4740011 0.5259989
1003 0.4927877 0.5072123
1004 0.4938562 0.5061438
```

```
lda.class[1:20]
```

```

[1] Up  Up  Up  Up  Up  Up  Up  Up  Up  Up  Up  Up  Down Up  Up  Up
[16] Up  Up  Down Up  Up
Levels: Down Up
```

```
sum(lda.pred$posterior[, 1] > .9)
```

```
[1] 0
```

```
qda.fit <- qda(Direction ~ Lag1 + Lag2, data = Smarket,
               subset = train)
qda.fit
```

Call:

```
qda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
```

Prior probabilities of groups:

	Down	Up
	0.491984	0.508016

Group means:

	Lag1	Lag2
Down	0.04279022	0.03389409
Up	-0.03954635	-0.03132544

```
qda.class <- predict(qda.fit, Smarket.2005)$class
table(qda.class, Direction.2005)
```

	Direction.2005	
qda.class	Down	Up
Down	30	20
Up	81	121

```
mean(qda.class == Direction.2005)
```

```
[1] 0.5992063
```

#Naive Bayes

```
nb.fit <- naiveBayes(Direction ~ Lag1 + Lag2, data = Smarket,
                     subset = train)
```

```
nb.fit
```

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = X, y = Y, laplace = laplace)
```

A-priori probabilities:

Y

	Down	Up
0.491984	0.508016	

Conditional probabilities:

		Lag1	
Y		[,1]	[,2]
	Down	0.04279022	1.227446
	Up	-0.03954635	1.231668

		Lag2	
Y		[,1]	[,2]
	Down	0.03389409	1.239191
	Up	-0.03132544	1.220765

```
mean(Lag1[train][Direction[train] == "Down"])
```

```
[1] 0.04279022
```

```
nb.class <- predict(nb.fit, Smarket.2005)
table(nb.class, Direction.2005)
```

		Direction.2005	
nb.class	Down	Up	
	Down	28	20
	Up	83	121

```
mean(nb.class == Direction.2005)
```

```
[1] 0.5912698
```

```
nb.preds <- predict(nb.fit, Smarket.2005, type = "raw")
nb.pred <- rep("Down", 252)
nb.pred[nb.preds[,2] > .5] <- "Up"
table(nb.pred, Direction.2005)
```

```
      Direction.2005
nb.pred Down  Up
Down    28  20
Up      83 121
```

```
#K-Nearest Neighbors
train.X <- cbind(Lag1, Lag2)[train, ]
test.X <- cbind(Lag1, Lag2)[!train, ]
train.Direction <- Direction[train]
```

```
set.seed(1)
knn.pred <- knn(train.X, test.X, train.Direction, k = 1)
table(knn.pred, Direction.2005)
```

```
      Direction.2005
knn.pred Down  Up
Down     43  58
Up       68  83
```

```
knn.pred <- knn(train.X, test.X, train.Direction, k = 3)
table(knn.pred, Direction.2005)
```

```
      Direction.2005
knn.pred Down  Up
Down     48  54
Up       63  87
```

```
mean(knn.pred == Direction.2005)
```

```
[1] 0.5357143
```

```
dim(Caravan)
```

```
[1] 5822    86
```

```
attach(Caravan)
summary(Purchase)
```

```
   No   Yes
5474  348
```

```
standardized.X <- scale(Caravan[, -86])#only drop Purchase
```

```
test <- 1:1000
train.X <- standardized.X[-test, ]
test.X <- standardized.X[test, ]
train.Y <- Purchase[-test]
test.Y <- Purchase[test]
```

```
set.seed(1)
knn.pred <- knn(train.X, test.X, train.Y, k = 1)
mean(test.Y != knn.pred)
```

```
[1] 0.118
```

```
mean(test.Y != "No")
```

```
[1] 0.059
```

```
table(knn.pred, test.Y)
```

```
      test.Y
knn.pred No Yes
   No   873  50
   Yes   68   9
```

```
knn.pred <- knn(train.X, test.X, train.Y, k = 3)
table(knn.pred, test.Y)
```

```
      test.Y
knn.pred No Yes
      No  920  54
      Yes   21   5
```


```
knn.pred <- knn(train.X, test.X, train.Y, k = 5)
table(knn.pred, test.Y)
```

```
      test.Y
knn.pred No Yes
      No  930  55
      Yes   11   4
```

#####exercise

13.

```
Weekly %>%
  ggplot(aes(x = Year, y = Volume)) +
  geom_point()
```



```
glm.fits <- glm(
  Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,
  data = Weekly,
  family = binomial
)
summary(glm.fits) #lag2
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
     Volume, family = binomial, data = Weekly)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.26686	0.08593	3.106	0.0019 **
Lag1	-0.04127	0.02641	-1.563	0.1181
Lag2	0.05844	0.02686	2.175	0.0296 *
Lag3	-0.01606	0.02666	-0.602	0.5469
Lag4	-0.02779	0.02646	-1.050	0.2937
Lag5	-0.01447	0.02638	-0.549	0.5833
Volume	-0.02274	0.03690	-0.616	0.5377

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2 on 1088 degrees of freedom  
Residual deviance: 1486.4 on 1082 degrees of freedom  
AIC: 1500.4

Number of Fisher Scoring iterations: 4

```
glm.probs <- predict(glm.fits, Weekly,  
                     type = "response")  
glm.pred <- rep("Down", 1089)  
glm.pred[glm.probs > .5] = "Up"  
table(glm.pred, Weekly$Direction)
```

```
glm.pred Down Up  
Down    54  48  
Up     430 557
```

```
attach(Weekly)
```

The following objects are masked from Smarket:

Direction, Lag1, Lag2, Lag3, Lag4, Lag5, Today, Volume, Year

```
train <- (Year < 2009)  
Weekly.test <- Weekly[!train, ]#test df  
Direction.test <- Direction[!train]  
  
glm.fits <- glm(  
  Direction ~ Lag2,  
  data = Weekly, family = binomial, subset = train  
)  
glm.probs <- predict(glm.fits, Weekly.test, type = "response")
```

#logistic



```
glm.pred <- rep("Down", 104)
glm.pred[glm.probs > .5] <- "Up"
table(glm.pred, Direction.test)
```

```

      Direction.test
glm.pred Down Up
Down      9  5
Up       34 56
```

```
mean(glm.pred == Direction.test)
```

```
[1] 0.625
```

#LDA

```
lda.fit <- lda(Direction ~ Lag2,
               data = Weekly,
               subset = train)

lda.pred <- predict(lda.fit, Weekly.test)
lda.class <- lda.pred$class
table(lda.class, Direction.test)
```

```

      Direction.test
lda.class Down Up
Down      9  5
Up       34 56
```

#QDA

```
qda.fit <- qda(Direction ~ Lag2,
               data = Weekly,
               subset = train)

qda.pred <- predict(qda.fit, Weekly.test)
qda.class <- qda.pred$class
table(qda.class, Direction.test)
```

```

      Direction.test
qda.class Down Up
      Down    0  0
      Up     43 61

```

```
mean(qda.class == Direction.test)
```

```
[1] 0.5865385
```

```
###KNN with K = 1
```

```

train.X <- matrix(Weekly$Lag2[train], ncol = 1)
test.X <- matrix(Weekly$Lag2[!train], ncol = 1)
train.Direction <- Direction[train]

set.seed(1)
knn.pred <- knn(train.X, test.X, train.Direction, k = 1)
table(knn.pred, Direction.test)

```

```

      Direction.test
knn.pred Down Up
      Down   21 30
      Up    22 31

```

```
#h
```

```

nb.fit <- naiveBayes(Direction ~ Lag2, data = Weekly,
                      subset = train)
nb.class <- predict(nb.fit, Weekly.test)
table(nb.class, Direction.test)

```

```

      Direction.test
nb.class Down Up
      Down    0  0
      Up     43 61

```

```
mean(nb.class == Direction.test)
```

```
[1] 0.5865385
```

```
#j Experiment
```

```
for (K in 1:5) {  
  train.X <- matrix(Weekly$Lag2[train], ncol = 1)  
  test.X <- matrix(Weekly$Lag2[!train], ncol = 1)  
  train.Direction <- Direction[train]  
  set.seed(1)  
  knn.pred <- knn(train.X, test.X, train.Direction, k = K)  
  C <- table(knn.pred, Direction.test)  
  
  if ("Up" %in% rownames(C)) {  
    pred <- sum(C["Up",])  
    did_increase <- C["Up", "Up"]  
    accuracy <- did_increase / pred # Ensure 'pred' is not zero before division  
  
    # Printing results  
    cat(sprintf("K=%d: # predicted to be up: %2d, # who did increase %d, accuracy %.1f%%\n",  
                K, pred, did_increase, accuracy * 100))  
  }  
}
```

```
K=1: # predicted to be up: 53, # who did increase 31, accuracy 58.5%  
K=2: # predicted to be up: 58, # who did increase 34, accuracy 58.6%  
K=3: # predicted to be up: 68, # who did increase 41, accuracy 60.3%  
K=4: # predicted to be up: 67, # who did increase 44, accuracy 65.7%  
K=5: # predicted to be up: 67, # who did increase 40, accuracy 59.7%
```

```
#14
```

```
Auto <- read.csv("Auto.csv")  
Auto$horsepower <- as.numeric(Auto$horsepower)
```

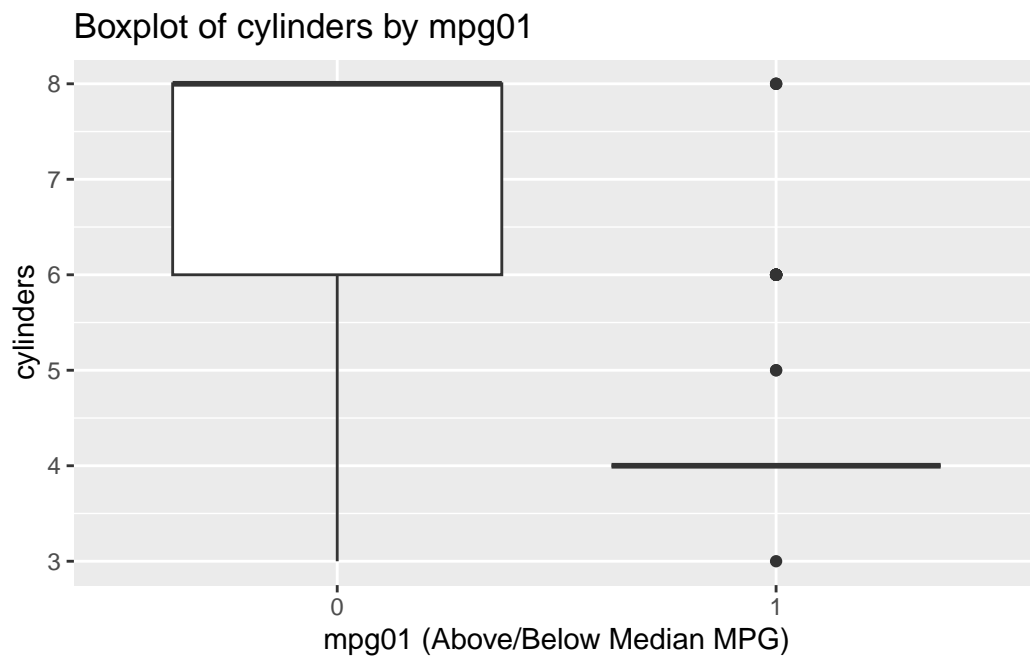
```
Warning: NAs introduced by coercion
```

```
Auto <- Auto[!is.na(Auto$horsepower), ]  
  
Auto$mpg01 <- if_else(Auto$mpg > median(Auto$mpg), 1, 0)
```


```
vars <- c('cylinders', 'displacement', 'horsepower', 'weight', 'acceleration')

# Loop through the variable names
for (i in vars) {
  # Use ggplot to create a box plot
  p <- ggplot(Auto, aes(x = factor(mpg01), y = get(i))) +
    geom_boxplot() +
    labs(title = paste('Boxplot of', i, 'by mpg01'),
         x = 'mpg01 (Above/Below Median MPG)',
         y = i)


  # Print the plot
  print(p)
}
```




Boxplot of displacement by mpg01



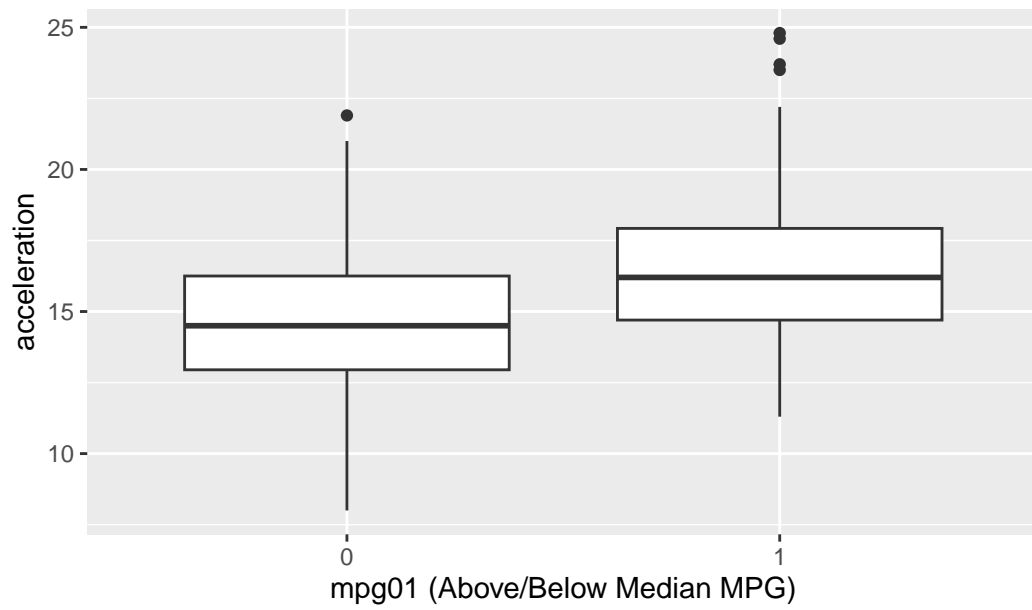
Boxplot of horsepower by mpg01



Boxplot of weight by mpg01



Boxplot of acceleration by mpg01



```
#'displacement' 'weight'
set.seed(1)
Auto_df <- Auto[, c('displacement', 'weight', 'mpg01')]
test <- sample(nrow(Auto_df), 100)

train.X <- Auto_df[-test, -3]
test.X <- Auto_df[test, -3]
train.df <- Auto_df[-test, ]

train.Y <- Auto_df[-test, 'mpg01']
test.Y <- Auto_df[test, 'mpg01']
```

#d LDA

```
lda.fit <- lda(mpg01 ~ displacement + weight,
               data = train.df)

lda.pred <- predict(lda.fit, test.X)
lda.class <- lda.pred$class
table(lda.class, test.Y)
```

```
      test.Y
lda.class 0  1
      0 43  4
      1  7 46
```

```
mean(lda.class == test.Y)
```

```
[1] 0.89
```

#QDA

```
qda.fit <- qda(mpg01 ~ displacement + weight,
               data = train.df)

qda.pred <- predict(qda.fit, test.X)
qda.class <- qda.pred$class
table(qda.class, test.Y)
```

```

      test.Y
qda.class 0  1
      0 47  5
      1  3 45

```

```
mean(qda.class == test.Y)
```

```
[1] 0.92
```

```
#logistic regression
```

```

glm.fits <- glm(
  mpg01 ~ displacement + weight,
  data = train.df, family = binomial
)

glm.probs <- predict(glm.fits, test.X,
  type = "response")

```

```

glm.pred <- rep("0", 100)
glm.pred[glm.probs > .5] <- "1"
table(glm.pred, test.Y)

```

```

      test.Y
glm.pred 0  1
      0 46  5
      1  4 45

```

```
mean(glm.pred == test.Y)
```

```
[1] 0.91
```



## naive Bayes

```
nb.fit <- naiveBayes(mpg01 ~ displacement + weight,  
                     data = train.df)  
nb.class <- predict(nb.fit, test.X)  
table(nb.class, test.Y)
```

```
      test.Y  
nb.class 0  1  
      0 45  4  
      1  5 46
```

```
mean(nb.class == test.Y)
```

```
[1] 0.91
```

```
knn.pred <- knn(train.X, test.X, train.Y, k = 1)  
C <- table(Predicted = knn.pred, Actual = test.Y)  
print(C)
```

```
      Actual  
Predicted 0  1  
      0 43  8  
      1  7 42
```

```
for (K in 1:5) {  
  set.seed(1)  
  knn.pred <- knn(train.X, test.X, train.Y, k = K)  
  C <- table(Predicted = knn.pred, Actual = test.Y)
```

```
  if ("1" %in% rownames(C)) {  
    pred <- sum(C["1",])  
    did_higher <- C["1", "1"]  
    if (pred > 0) {  
      accuracy <- did_higher / pred
```

```
    # Printing results
```

```

        cat(sprintf("K=%d: # predicted to be higher than median: %2d, # who did higher than
                    K, pred, did_higher, accuracy * 100))
    }
}
}

```

```

K=1: # predicted to be higher than median: 49, # who did higher than median 42, accuracy 85.7
K=2: # predicted to be higher than median: 50, # who did higher than median 44, accuracy 88.
K=3: # predicted to be higher than median: 50, # who did higher than median 44, accuracy 88.
K=4: # predicted to be higher than median: 50, # who did higher than median 45, accuracy 90.
K=5: # predicted to be higher than median: 52, # who did higher than median 46, accuracy 88.

```

```
#15
```

```

Power <- function() {
  result <- 2^3
  print(result)
}
Power()

```

```
[1] 8
```

```

Power2 <- function(x, a) {
  result <- x^a
  print(result)
}
Power2(3, 8)

```

```
[1] 6561
```

```
Power2(10, 3)
```

```
[1] 1000
```

```
Power2(8, 17)
```

```
[1] 2.2518e+15
```

```
Power2(13, 13)
```

```
[1] 3.028751e+14
```

```
Power3 <- function(x, a) {  
  result <- x^a  
  return(result)  
}
```


```
value1 <- Power3(2, 3)
```

```
x_values <- 1:10  
y_values <- x_values^2
```

```
df <- data.frame(x_values, y_values)
```


```
df %>%  
  ggplot( aes(x = x_values, y = y_values)) +  
  geom_point() +  
  geom_smooth() +  
  ggtitle("Plot of  $f(x) = x^2$ ") +  
  xlab("x values") +  
  ylab(" $f(x) = x^2$  values")
```

`geom\_smooth()` using method = 'loess' and formula = 'y ~ x'



```
PlotPower <- function(x_values, a) {  
  y_values <- x_values^a  
  
  df <- data.frame(x_values, y_values)  
  
  df %>%  
    ggplot( aes(x = x_values, y = y_values)) +  
    geom_point() +  
    geom_smooth() +  
    ggtitle("Plot of f(x) = x^2") +  
    xlab("x values") +  
    ylab("f(x) = x^2 values")  
}  
  
PlotPower(1:10, 3)
```

`geom\_smooth()` using method = 'loess' and formula = 'y ~ x'



#16.

```
Boston$crim01 <- if_else(Boston$crim > median(Boston$crim), 1, 0)


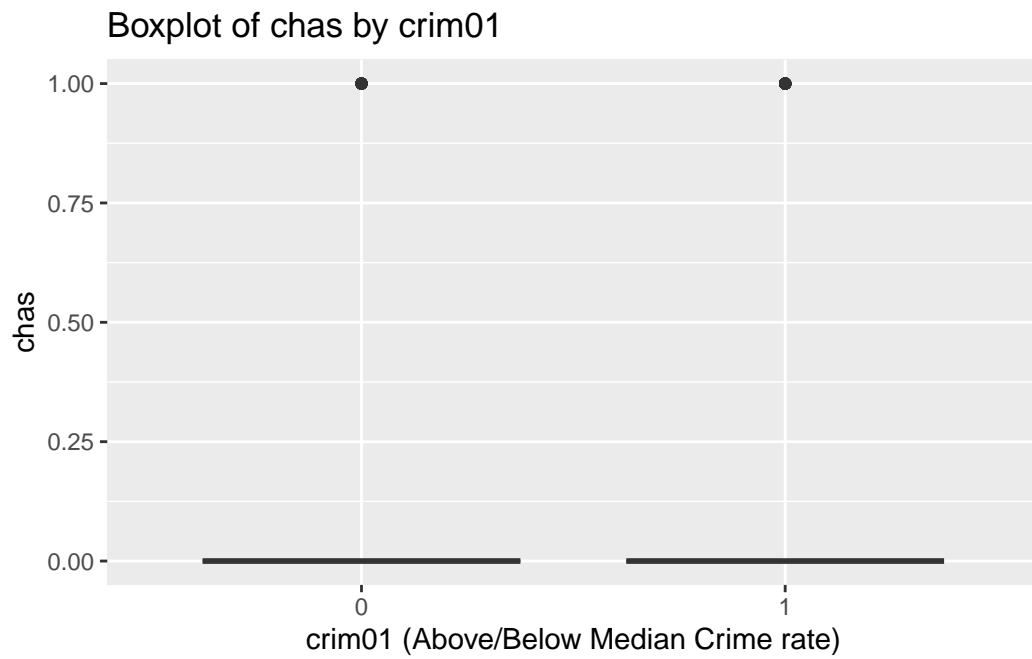
vars <- names(Boston)[which(names(Boston) == "zn"):which(names(Boston) == "medv")]

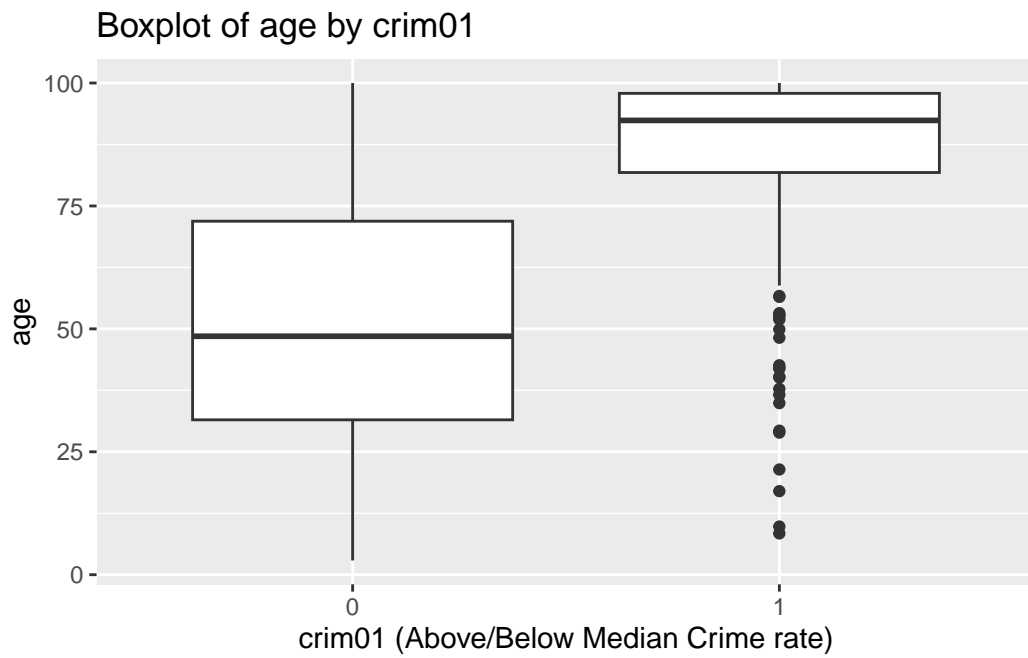

# Loop through the variable names
for (i in vars) {
  # Use ggplot to create a box plot
  p <- ggplot(Boston, aes(x = factor(crim01), y = get(i))) +
    geom_boxplot() +
    labs(title = paste('Boxplot of', i, 'by crim01'),
         x = 'crim01 (Above/Below Median Crime rate)',
         y = i)

  # Print the plot
  print(p)
}
```

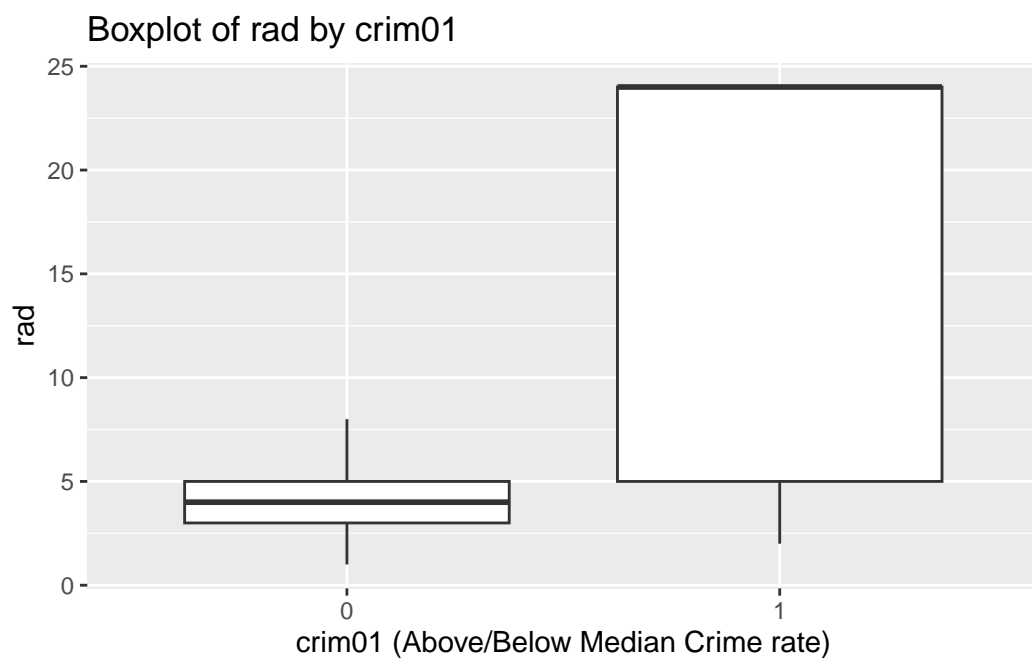
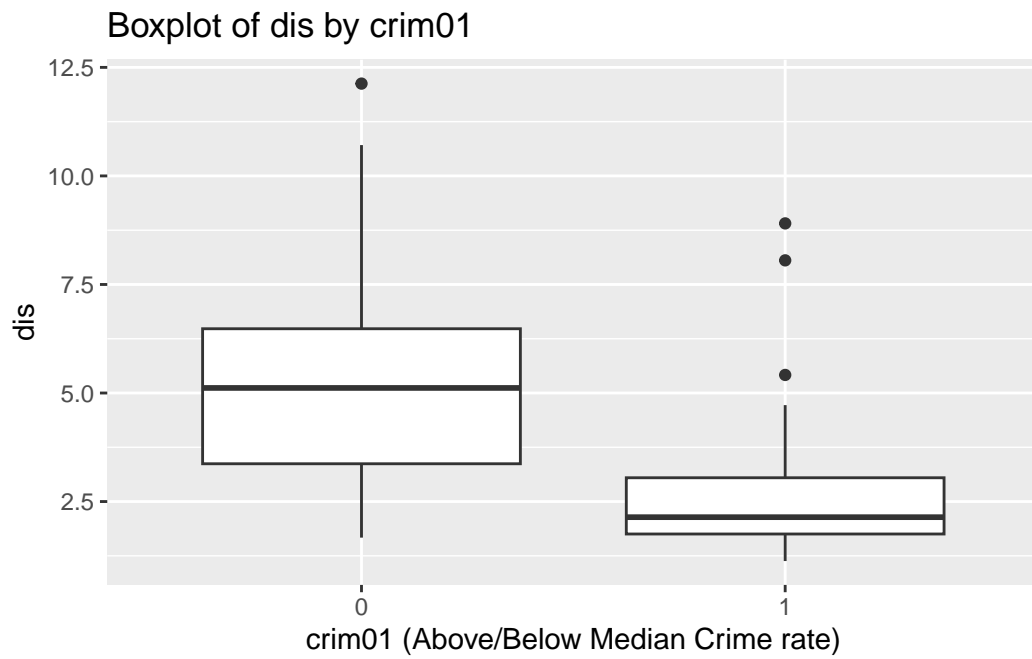
A boxplot comparing the distribution of the 'indus' variable across two categories. The y-axis is labeled 'indus' and ranges from 0 to 20. The left box (category 1) has a median around 5.5, with the box spanning from approximately 3 to 8.5. The right box (category 2) has a median around 15, with the box spanning from approximately 10 to 18.5. Both categories show whiskers extending to the minimum and maximum values, with several outliers plotted as black dots above the upper whisker for both groups.

30












Boxplot of tax by crim01




Boxplot of ptratio by crim01

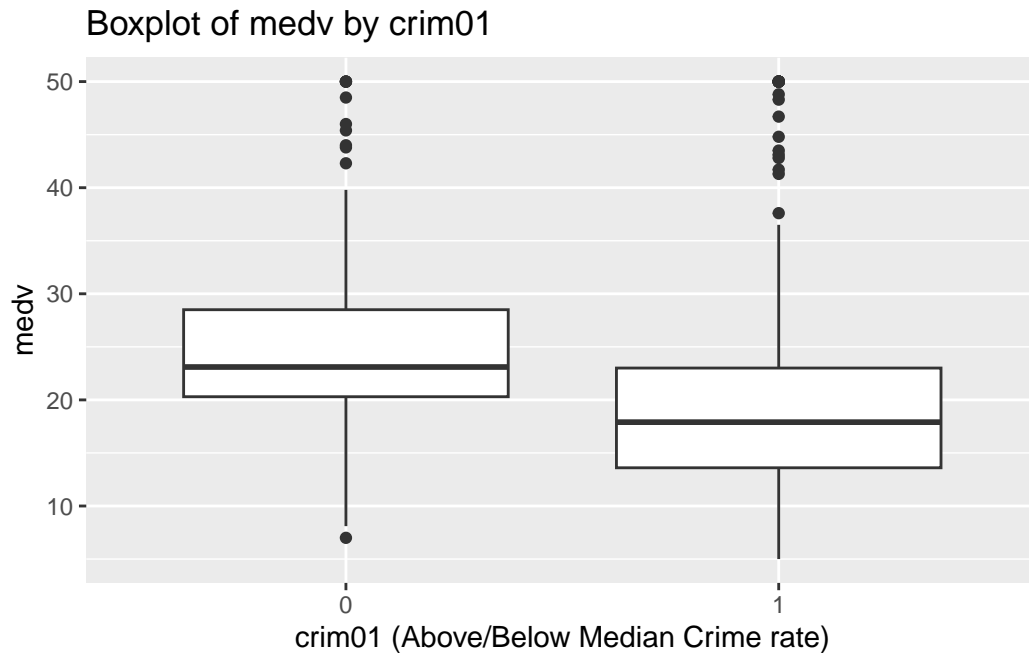


Boxplot of black by crim01



Boxplot of lstat by crim01





```
#indus,nox,age,dis,rad
```

```
set.seed(1)
Boston_df <- Boston[, c('indus', 'nox', 'age', 'dis', 'rad', 'crim01')]
test <- sample(nrow(Boston_df), 200)

train.X <- Boston_df[-test, -6]
test.X <- Boston_df[test, -6]
train.df <- Boston_df[-test, ]

train.Y <- Boston_df[-test, 'crim01']
test.Y <- Boston_df[test, 'crim01']
```

```
#logistic
```

```
glm.fits <- glm(
  crim01 ~ indus + nox + age + dis + rad,
  data = train.df, family = binomial
)

glm.probs <- predict(glm.fits, test.X,
```

```

                                type = "response")

glm.pred <- rep("0", 200)
glm.pred[glm.probs > .5] <- "1"
table(glm.pred, test.Y)

      test.Y
glm.pred  0  1
      0 97 16
      1  9 78

mean(glm.pred == test.Y)

[1] 0.875

#LDA

lda.fit <- lda(crim01 ~ indus + nox + age +dis + rad,
               data = train.df)

lda.pred <- predict(lda.fit, test.X)
lda.class <- lda.pred$class
table(lda.class, test.Y)

      test.Y
lda.class  0  1
      0 95 23
      1 11 71

mean(lda.class == test.Y)

[1] 0.83

#qda

```

```
qda.fit <- qda(crim01 ~ indus + nox + age +dis + rad,
               data = train.df)
```

```
qda.pred <- predict(qda.fit, test.X)
qda.class <- qda.pred$class
table(qda.class, test.Y)
```

```
      test.Y
qda.class 0  1
      0 99 18
      1  7 76
```

```
mean(qda.class == test.Y)
```

```
[1] 0.875
```

```
#naiveBayes
```

```
nb.fit <- naiveBayes(crim01 ~ indus + nox + age +dis + rad,
                     data = train.df)
nb.fit
```

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = X, y = Y, laplace = laplace)
```

A-priori probabilities:

```
Y
      0      1
0.4803922 0.5196078
```

Conditional probabilities:

```
      indus
Y      [,1]      [,2]
0  6.681088 5.289341
1 15.132893 5.550083
```

```

      nox
Y      [,1]      [,2]
0 0.4681619 0.05532716
1 0.6397610 0.09878089

```

```

      age
Y      [,1]      [,2]
0 51.71361 25.69922
1 86.19434 17.05572

```

```

      dis
Y      [,1]      [,2]
0 5.188369 2.099384
1 2.456703 1.048184

```

```

      rad
Y      [,1]      [,2]
0  4.163265 1.613474
1 14.572327 9.565536

```

```

nb.class <- predict(nb.fit, test.X)
table(nb.class, test.Y)

```

```

      test.Y
nb.class 0  1
0  92 20
1  14 74

```