

# Untitled

```
install.packages("e1071")
```

Installing package into '/cloud/lib/x86\_64-pc-linux-gnu-library/4.4'  
(as 'lib' is unspecified)


```
library(e1071)
```

```
#####labs
```

## 9.6.1 Support Vector Classifier

#checking whether the classes are linearly separable

```
set.seed(1)
x <- matrix(rnorm(20 * 2), ncol = 2)
y <- c(rep(-1, 10), rep(1, 10))
x[y == 1, ] <- x[y == 1, ] + 1
plot(x, col = (3 - y))
```




##fit the classifier

```
dat <- data.frame(x = x, y = as.factor(y))
svmfit <- svm(y ~ ., data = dat, kernel = "linear",
              cost = 10, scale = FALSE)
```

```
plot(svmfit, dat)
```

### SVM classification plot



```
svmfit$index
```

```
[1] 1 2 5 7 14 16 17
```

#there are seven support vectors

```
summary(svmfit)
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 10, scale = FALSE)
```

Parameters:

```
  SVM-Type:  C-classification
 SVM-Kernel:  linear
      cost:  10
```

Number of Support Vectors: 7

```
( 4 3 )
```


Number of Classes: 2

Levels:

```
-1 1
```

```
svmfit <- svm(y ~ ., data = dat, kernel = "linear",
              cost = 0.1, scale = FALSE)
plot(svmfit, dat)
```

## SVM classification plot



```
svmfit$index
```

```
[1] 1 2 3 4 5 7 9 10 12 13 14 15 16 17 18 20
```

#Now that a smaller value of the cost parameter is being used, we obtain a larger number of support vectors, because the margin is now wider.

#perform cross validation

```
set.seed(1)
tune.out <- tune(svm, y ~ ., data = dat, kernel = "linear",
  ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
```

```
summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
cost
0.1
```

- best performance: 0.05

- Detailed performance results:

	cost	error	dispersion
1	1e-03	0.55	0.4377975
2	1e-02	0.55	0.4377975
3	1e-01	0.05	0.1581139
4	1e+00	0.15	0.2415229
5	5e+00	0.15	0.2415229
6	1e+01	0.15	0.2415229
7	1e+02	0.15	0.2415229

#cost = 0.1 results in the lowest cross-validation error rate

```
bestmod <- tune.out$best.model  
summary(bestmod)
```

Call:

```
best.tune(METHOD = svm, train.x = y ~ ., data = dat, ranges = list(cost = c(0.001,  
  0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
```

Parameters:

```
SVM-Type:  C-classification  
SVM-Kernel: linear  
cost:  0.1
```

Number of Support Vectors: 16

( 8 8 )

Number of Classes: 2

Levels:

-1 1

#generating a test data set.

```
xtest <- matrix(rnorm(20 * 2), ncol = 2)
ytest <- sample(c(-1, 1), 20, rep = TRUE)
xtest[ytest == 1, ] <- xtest[ytest == 1, ] + 1
testdat <- data.frame(x = xtest, y = as.factor(ytest))
```

**predict the class labels of these test observations**

```
ypred <- predict(bestmod, testdat)
table(predict = ypred, truth = testdat$y)
```

```
      truth
predict -1 1
      -1  9 1
       1  2 8
```

**What if we had instead used  $\text{cost} = 0.01$ ?**


```
svmfit <- svm(y ~ ., data = dat, kernel = "linear",
              cost = .01, scale = FALSE)
ypred <- predict(svmfit, testdat)
table(predict = ypred, truth = testdat$y)
```

```
      truth
predict -1 1
      -1 11 6
       1  0 3
```

**three additional observations are misclassified**

**further separate the two classes in our simulated data so that they are linearly separable**

```
x[y == 1, ] <- x[y == 1, ] + 0.5
plot(x, col = (y + 5) / 2, pch = 19)
```



```
dat <- data.frame(x = x, y = as.factor(y))
svmfit <- svm(y ~ ., data = dat, kernel = "linear",
              cost = 1e5)
summary(svmfit)
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 1e+05)
```

Parameters:

```
SVM-Type:  C-classification
SVM-Kernel: linear
cost:      1e+05
```

Number of Support Vectors: 3

```
( 1 2 )
```

Number of Classes: 2

Levels:

```
-1 1
```

#only three support vectors were used

#It seems likely that this model will perform poorly on test data.

```
svmfit <- svm(y ~ ., data = dat, kernel = "linear", cost = 1)
summary(svmfit)
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 1)
```

Parameters:

SVM-Type: C-classification

SVM-Kernel: linear

cost: 1

Number of Support Vectors: 7

( 4 3 )

Number of Classes: 2


Levels:

-1 1

```
plot(svmfit, dat)
```




## SVM classification plot



#9.6.2 Support Vector Machine

```
set.seed(1)
x <- matrix(rnorm(200 * 2), ncol = 2)
x[1:100, ] <- x[1:100, ] + 2
x[101:150, ] <- x[101:150, ] - 2
y <- c(rep(1, 150), rep(2, 50))
dat <- data.frame(x = x, y = as.factor(y))
```


```
plot(x, col = y)
```



fit the training data using the `svm()` function with a radial kernel and `gamma = 1`:

```
train <- sample(200, 100)
svmfit <- svm(y ~ ., data = dat[train, ], kernel = "radial",
             gamma = 1, cost = 1)
plot(svmfit, dat[train, ])
```

**SVM classification plot**



```
summary(svmfit)
```

Call:

```
svm(formula = y ~ ., data = dat[train, ], kernel = "radial", gamma = 1,
    cost = 1)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1
```

Number of Support Vectors: 31

( 16 15 )

Number of Classes: 2

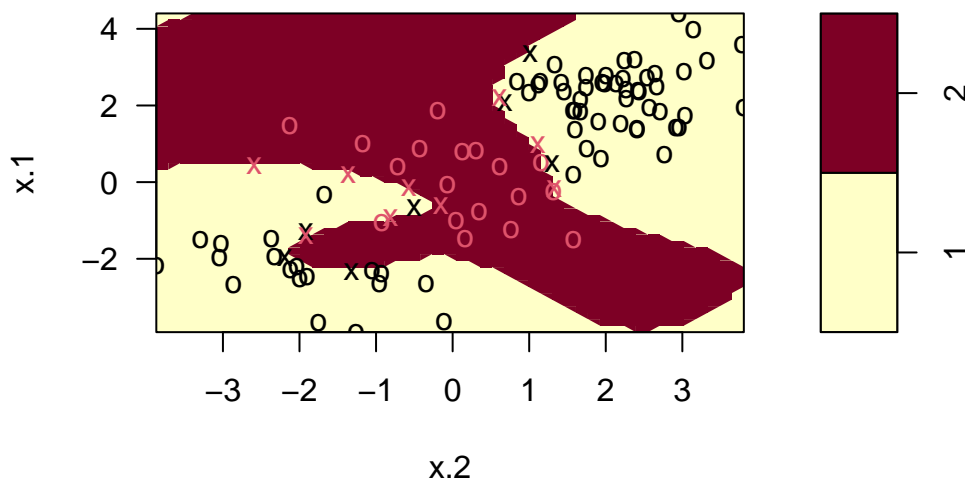
Levels:

1 2

seems to be at risk of overfitting the data.

```
svmfit <- svm(y ~ ., data = dat[train, ], kernel = "radial",  
             gamma = 1, cost = 1e5)  
plot(svmfit, dat[train, ])
```

**SVM classification plot**



#perform cross-validation using tune() to select the best choice of  $\gamma$  and cost for an SVM with a radial kernel

```
set.seed(1)  
tune.out <- tune(svm, y ~ ., data = dat[train, ],  
               kernel = "radial",  
               ranges = list(cost = c(0.1, 1, 10, 100, 1000),  
                             gamma = c(0.5, 1, 2, 3, 4)  
               ) )  
summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
cost gamma
1 0.5
```

- best performance: 0.07

- Detailed performance results:

	cost	gamma	error	dispersion
1	1e-01	0.5	0.26	0.15776213
2	1e+00	0.5	0.07	0.08232726
3	1e+01	0.5	0.07	0.08232726
4	1e+02	0.5	0.14	0.15055453
5	1e+03	0.5	0.11	0.07378648
6	1e-01	1.0	0.22	0.16193277
7	1e+00	1.0	0.07	0.08232726
8	1e+01	1.0	0.09	0.07378648
9	1e+02	1.0	0.12	0.12292726
10	1e+03	1.0	0.11	0.11005049
11	1e-01	2.0	0.27	0.15670212
12	1e+00	2.0	0.07	0.08232726
13	1e+01	2.0	0.11	0.07378648
14	1e+02	2.0	0.12	0.13165612
15	1e+03	2.0	0.16	0.13498971
16	1e-01	3.0	0.27	0.15670212
17	1e+00	3.0	0.07	0.08232726
18	1e+01	3.0	0.08	0.07888106
19	1e+02	3.0	0.13	0.14181365
20	1e+03	3.0	0.15	0.13540064
21	1e-01	4.0	0.27	0.15670212
22	1e+00	4.0	0.07	0.08232726
23	1e+01	4.0	0.09	0.07378648
24	1e+02	4.0	0.13	0.14181365
25	1e+03	4.0	0.15	0.13540064

```
table(
  true = dat[-train, "y"],
  pred = predict(
```

```
tune.out$best.model, newdata = dat[-train, ]
) )
```

```
      pred
true  1  2
     1 67 10
     2  2 21
```


#12 % of test observations are misclassified by this SVM

### #9.6.3 ROC Curves

```
library(ROCR)
rocplot <- function(pred, truth, ...) {
  predob <- prediction(pred, truth)
  perf <- performance(predob, "tpr", "fpr")
  plot(perf, ...)
}
```

```
svmfit.opt <- svm(y ~ ., data = dat[train, ],
                 kernel = "radial", gamma = 2, cost = 1,
                 decision.values = T)
fitted <- attributes(
  predict(svmfit.opt, dat[train, ], decision.values = TRUE)
)$decision.values
```

```
par(mfrow = c(1, 2))
rocplot(-fitted, dat[train, "y"], main = "Training Data")
```



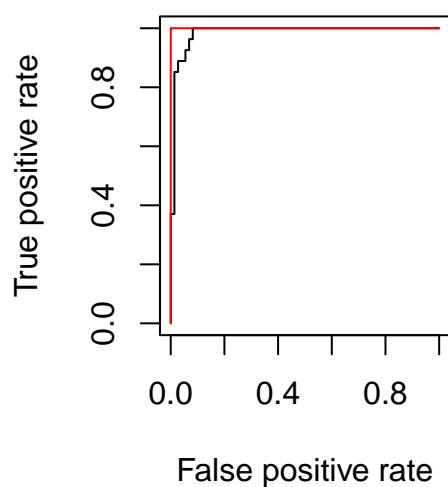
#SVM appears to be producing accurate predictions. By increasing `gamma` we can produce a more flexible fit and generate further improvements in accuracy

```
par(mfrow = c(1, 2))
rocplot(-fitted, dat[train, "y"], main = "Training Data")
svmfit.flex <- svm(y ~ ., data = dat[train, ],
                  kernel = "radial", gamma = 50, cost = 1,
                  decision.values = T)

fitted <- attributes(
  predict(svmfit.flex, dat[train, ], decision.values = T)
)$decision.values

rocplot(-fitted, dat[train, "y"], add = T, col = "red")
```


## Training Data



#We are really more interested in the level of prediction accuracy on the test data. #model with  $\gamma = 2$  appears to provide the most accurate results.


```
fitted <- attributes(  
  predict(svmfit.opt, dat[-train, ], decision.values = T)  
)$decision.values  
  
rocplot(-fitted, dat[-train, "y"], main = "Test Data")  
fitted <- attributes(  
  predict(svmfit.flex, dat[-train, ], decision.values = T)  
)$decision.values  
  
rocplot(-fitted, dat[-train, "y"], add = T, col = "red")
```

## Test Data



### #9.6.4 SVM with Multiple Classes


```
set.seed(1)
x <- rbind(x, matrix(rnorm(50 * 2), ncol = 2)) #entries for 50 observations, each with two f
y <- c(y, rep(0, 50)) #creates a vector of 50 zeros
x[y == 0, 2] <- x[y == 0, 2] + 2
dat <- data.frame(x = x, y = as.factor(y))
par(mfrow = c(1, 1))
plot(x, col = (y + 1))
```





```
svmfit <- svm(y ~ ., data = dat, kernel = "radial",
             cost = 10, gamma = 1)
plot(svmfit, dat)
```

## SVM classification plot



### #9.6.5 Application to Gene Expression Data

```
library(ISLR2)
names(Khan)
```

```
[1] "xtrain" "xtest"  "ytrain" "ytest"
```

```
dim(Khan$xtrain)
```

```
[1] 63 2308
```

```
dim(Khan$xtest)
```

```
[1] 20 2308
```

```
length(Khan$ytrain)
```

```
[1] 63
```

```
length(Khan$ytest)
```

```
[1] 20
```

```
table(Khan$ytrain)
```

```
1  2  3  4  
8 23 12 20
```

```
table(Khan$ytest)
```

```
1 2 3 4  
3 6 6 5
```

```
dat <- data.frame(  
  x = Khan$xtrain,  
  y = as.factor(Khan$ytrain)  
)  
  
out <- svm(y ~ ., data = dat, kernel = "linear",  
           cost = 10)  
  
summary(out)
```

Call:

```
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 10)
```

Parameters:

```
SVM-Type:  C-classification  
SVM-Kernel: linear  
cost:      10
```

Number of Support Vectors: 58

```
( 20 20 11 7 )
```

Number of Classes: 4

Levels:

1 2 3 4

```
table(out$fitted, dat$y)
```

	1	2	3	4
1	8	0	0	0
2	0	23	0	0
3	0	0	12	0
4	0	0	0	20


```
dat.te <- data.frame(  
  x = Khan$xtest,  
  y = as.factor(Khan$ytest))  
  
pred.te <- predict(out, newdata = dat.te)  
table(pred.te, dat.te$y)
```

pred.te	1	2	3	4
1	3	0	0	0
2	0	6	2	0
3	0	0	4	0
4	0	0	0	5

#9.7 Applied Exercises #4.

```
set.seed(1)  
x <- matrix(rnorm(100 * 2), ncol = 2)  
x[1:20, ] <- x[1:20, ] + 1  
x[20:40, ] <- x[20:40, ] - 2  
y <- c(rep(1, 50), rep(2, 50))  
dat <- data.frame(x = x, y = as.factor(y))
```

```
plot(x, col = y)
```



#a polynomial kernel

```
train <- sample(100, 50)
svmfit <- svm(y ~ ., data = dat[train, ], kernel = "poly",
              degree = 2)


set.seed(1)
tune.out <- tune(svm, y ~ ., data = dat[train, ],
                 kernel = "poly",
                 ranges = list(degree = c(1, 2, 3, 4)
) )
summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation
- best parameters:  
degree  
4
- best performance: 0.22
- Detailed performance results:

	degree	error	dispersion
1	1	0.26	0.2836273
2	2	0.24	0.2065591
3	3	0.36	0.2270585
4	4	0.22	0.1988858

```
plot(svmfit, dat)
```



```
table(
  true = dat[-train, "y"],
  pred = predict(tune.out$best.model, newdata = dat[-train, ]) )
```

	pred	
true	1	2
1	18	10
2	4	18

#a radial kernel

```
svmfit <- svm(y ~ ., data = dat[train, ], kernel = "radial",
              gamma = 1)

set.seed(1)
tune.out <- tune(svm, y ~ ., data = dat[train, ],
```

```

        kernel = "radial",
        ranges = list(cost = c(0.1, 1, 10, 100, 1000),
                      gamma = c(0.5, 1, 2, 3, 4)
    ) )
summary(tune.out)

```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```

cost gamma
1      0.5

```


- best performance: 0.24

- Detailed performance results:

	cost	gamma	error	dispersion
1	1e-01	0.5	0.44	0.2796824
2	1e+00	0.5	0.24	0.2458545
3	1e+01	0.5	0.26	0.2319004
4	1e+02	0.5	0.30	0.2538591
5	1e+03	0.5	0.36	0.2458545
6	1e-01	1.0	0.44	0.2796824
7	1e+00	1.0	0.24	0.2458545
8	1e+01	1.0	0.30	0.2538591
9	1e+02	1.0	0.32	0.2347576
10	1e+03	1.0	0.38	0.2740641
11	1e-01	2.0	0.44	0.2796824
12	1e+00	2.0	0.26	0.2319004
13	1e+01	2.0	0.32	0.2347576
14	1e+02	2.0	0.40	0.2981424
15	1e+03	2.0	0.48	0.2347576
16	1e-01	3.0	0.44	0.2796824
17	1e+00	3.0	0.30	0.2160247
18	1e+01	3.0	0.36	0.2796824
19	1e+02	3.0	0.42	0.2740641
20	1e+03	3.0	0.46	0.2836273
21	1e-01	4.0	0.44	0.2796824
22	1e+00	4.0	0.30	0.2357023
23	1e+01	4.0	0.38	0.3047768

```
24 1e+02 4.0 0.42 0.2898275
25 1e+03 4.0 0.46 0.2503331
```

```
plot(svmfit, dat)
```



```
table(
  true = dat[-train, "y"],
  pred = predict(tune.out$best.model, newdata = dat[-train, ]) )
```


```
      pred
true  1  2
  1 17 11
  2  3 19
```

**techniques performs almost the same**

#5.

```
set.seed(12)
x1 <- runif(500) - 0.5
x2 <- runif(500) - 0.5
y <- 1 * (x1^2 - x2^2 > 0)
```

```
par(mfrow = c(1, 1))
plot(x1, x2, col = (y+1) , pch = 3-y)
```



#c) Fit a logistic regression model to the data, using X1 and X2 as predictors

```
glm.fits <- glm(y ~ x1 + x2, family = binomial)
summary(glm.fits)
```

Call:

```
glm(formula = y ~ x1 + x2, family = binomial)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.04927	0.08978	0.549	0.583
x1	-0.23002	0.31534	-0.729	0.466
x2	0.51072	0.31560	1.618	0.106

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 692.86 on 499 degrees of freedom  
 Residual deviance: 689.58 on 497 degrees of freedom  
 AIC: 695.58

Number of Fisher Scoring iterations: 3

#d) Apply this model to the training data



```
glm.probs <- predict(glm.fits, type = "response")
glm.pred <- rep(0, 500)
glm.pred[glm.probs > .5] = 1
table(glm.pred, y)
```

```
      y
glm.pred  0   1
0 125  73
1 119 183
```

#e) Now fit a logistic regression model to the data using non-linear functions of X1 and X2 as predictors

```
glm.fits1 <- glm(y ~ x1 + I(x1^2) + x2, family = binomial)
summary(glm.fits1)
```

Call:

```
glm(formula = y ~ x1 + I(x1^2) + x2, family = binomial)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.3955	0.1661	-8.404	<2e-16 ***
x1	-0.4005	0.4036	-0.992	0.3211
I(x1^2)	19.1102	1.8827	10.151	<2e-16 ***
x2	0.8015	0.3696	2.169	0.0301 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 692.86 on 499 degrees of freedom  
 Residual deviance: 539.84 on 496 degrees of freedom  
 AIC: 547.84


Number of Fisher Scoring iterations: 4

#f) Apply this model to the training data in order to obtain a predicted class label for each training observation.

```
glm.probs1 <- predict(glm.fits1, type = "response")
glm.pred1 <- rep(0, 500)
glm.pred1[glm.probs1 > .5] = 1
table(glm.pred1, y)
```


```
      y
glm.pred1  0   1
0  189   78
1   55  178
```

```
data <- data.frame(x1 = x1, x2 = x2, y = y)
plot(data[glm.pred1 == 1, ]$x1, data[glm.pred1 == 1, ]$x2, col = (1 + 1), pch = (3 - 1), xlab = "X1", ylab = "X2")
points(data[glm.pred1 == 0, ]$x1, data[glm.pred1 == 0, ]$x2, col = (1 + 0), pch = (3 - 0))
```




#g) Fit a support vector classifier to the data with X1 and X2 as predictors.

```
data$y <- as.factor(data$y)
svmfit <- svm(y ~ x1 + x2, data = data, kernel = "linear",
              cost = 0.01, scale = FALSE)
preds <- predict(svmfit, data)
plot(data[preds == 1, ]$x1, data[preds == 1, ]$x2, col = (1 + 1), pch = (3 - 1), xlab = "X1", ylab = "X2")
points(data[preds == 0, ]$x1, data[preds == 0, ]$x2, col = (1 + 0), pch = (3 - 0))
```



#h) Fit a SVM using a non-linear kernel to the data

```
svmfit <- svm(y ~ x1 + x2, data = data, kernel = "radial",
             gamma = 1, cost = 1)
preds <- predict(svmfit, data)
plot(data[preds == 1, ]$x1, data[preds == 1, ]$x2, col = (1 + 1), pch = (3 - 1), xlab = "X1",
     points(data[preds == 0, ]$x1, data[preds == 0, ]$x2, col = (1 + 0), pch = (3 - 0)))
```




#i) Comment on your results. #fitting a SVM seems best

#6.

#a) Generate two-class data with  $p = 2$

```
set.seed(1)
x <- matrix(rnorm(100 * 2), ncol = 2)
y <- c(rep(-1, 50), rep(1, 50))
x[y == 1, ] <- x[y == 1, ] + 2.5
plot(x, col = (3 - y))
```



#b) Compute the cross-validation error rates for support vector classifiers with a range of cost values

```
dat <- data.frame(x = x, y = y)
set.seed(1)
tune.out <- tune(svm, y ~ ., data = dat, kernel = "linear",
                 ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
```

```
summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
cost
0.1
```

- best performance: 0.2168627

- Detailed performance results:

	cost	error	dispersion
1	1e-03	1.2669256	0.45202618
2	1e-02	0.2492452	0.05693377
3	1e-01	0.2168627	0.04430729
4	1e+00	0.2210400	0.04191765
5	5e+00	0.2209562	0.04155216
6	1e+01	0.2208784	0.04159409
7	1e+02	0.2208707	0.04156451

```
bestmod <- tune.out$best.model  
summary(bestmod)
```

Call:

```
best.tune(METHOD = svm, train.x = y ~ ., data = dat, ranges = list(cost = c(0.001,  
0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
```

Parameters:

```
SVM-Type:  eps-regression  
SVM-Kernel: linear  
cost:      0.1  
gamma:     0.5  
epsilon:   0.1
```

Number of Support Vectors: 87

#Number of Support Vectors: 87

```
preds <- predict(bestmod, dat)  
pred1 <- rep(0, 100)  
pred1[preds > .5] = 1  
table(pred1, y)
```

	y	
pred1	-1	1
	0	50 12
	1	0 38

#c) Generate an appropriate test data set

```
set.seed(12)
xtest <- matrix(rnorm(50 * 2), ncol = 2)
ytest <- c(rep(-1, 25), rep(1, 25))
xtest[ytest == 1, ] <- xtest[ytest == 1, ] + 2.5
testdat <- data.frame(x = xtest, y = as.factor(ytest))
```

```
costs <- c(0.01, 0.1, 1, 5, 10)
test.err <- rep(NA, length(costs))
data.test <- data.frame(x = xtest, z = as.factor(ytest))
for (i in 1:length(costs)) {
  svmfit <- svm(y ~ ., data = dat, kernel = "linear", cost = costs[i])
  pred <- predict(svmfit, testdat)
  pred1 <- rep(0, 50)
  pred1[pred > .5] = 1
  test.err[i] <- sum(pred1 != testdat$y)
}
data.frame(cost = costs, misclass = test.err)
```

	cost	misclass
1	0.01	35
2	0.10	31
3	1.00	31
4	5.00	31
5	10.00	31

#0.1 #7.

```
Auto <- na.omit(Auto)
Auto$mpg_b <- ifelse(Auto$mpg > median(Auto$mpg), 1, 0)
Auto$mpg_b <- as.factor(Auto$mpg_b )
```

```
tune.out <- tune(svm, mpg_b ~ .-mpg, data = Auto, kernel = "linear",
  ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost  
0.01

- best performance: 0.08935897

- Detailed performance results:

	cost	error	dispersion
1	1e-03	0.13025641	0.03349293
2	1e-02	0.08935897	0.03478144
3	1e-01	0.09185897	0.03843785
4	1e+00	0.09961538	0.02845044
5	5e+00	0.11750000	0.03895464
6	1e+01	0.11500000	0.04437783
7	1e+02	0.12519231	0.02882961

#cost = 1

```
bestmod <- tune.out$best.model  
summary(bestmod)
```

Call:

```
best.tune(METHOD = svm, train.x = mpg_b ~ . - mpg, data = Auto, ranges = list(cost = c(0.001  
0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
```

Parameters:

SVM-Type: C-classification  
SVM-Kernel: linear  
cost: 0.01

Number of Support Vectors: 168

( 84 84 )

Number of Classes: 2

Levels:

0 1

```
ypred <- predict(bestmod, Auto)

table(predict = ypred, truth = Auto$mpg_b)
```

```
      truth
predict 0   1
      0 170   8
      1  26 188
```

#c) Now repeat (b), this time using SVMs with radial and polynomial basis kernels,

```
tune.out <- tune(svm, mpg_b ~ .-mpg, data = Auto, kernel = "radial",
               ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100),
                             gamma = c(0.5, 1, 2, 3, 4)))
summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
cost gamma
  1      1
```

- best performance: 0.08147436

- Detailed performance results:

	cost	gamma	error	dispersion
1	1e-03	0.5	0.56358974	0.05130146
2	1e-02	0.5	0.56358974	0.05130146
3	1e-01	0.5	0.08923077	0.04214637
4	1e+00	0.5	0.08410256	0.03993932
5	5e+00	0.5	0.08653846	0.03367809
6	1e+01	0.5	0.08910256	0.03384038
7	1e+02	0.5	0.09166667	0.02914298
8	1e-03	1.0	0.56358974	0.05130146
9	1e-02	1.0	0.56358974	0.05130146
10	1e-01	1.0	0.56358974	0.05130146



11	1e+00	1.0	0.08147436	0.03917716
12	5e+00	1.0	0.08910256	0.03979296
13	1e+01	1.0	0.08910256	0.03979296
14	1e+02	1.0	0.09166667	0.04154431
15	1e-03	2.0	0.56358974	0.05130146
16	1e-02	2.0	0.56358974	0.05130146
17	1e-01	2.0	0.56358974	0.05130146
18	1e+00	2.0	0.12987179	0.07213189
19	5e+00	2.0	0.13243590	0.06943307
20	1e+01	2.0	0.13243590	0.06943307
21	1e+02	2.0	0.13243590	0.06943307
22	1e-03	3.0	0.56358974	0.05130146
23	1e-02	3.0	0.56358974	0.05130146
24	1e-01	3.0	0.56358974	0.05130146
25	1e+00	3.0	0.36455128	0.18145896
26	5e+00	3.0	0.36711538	0.17204930
27	1e+01	3.0	0.36711538	0.17204930
28	1e+02	3.0	0.36711538	0.17204930
29	1e-03	4.0	0.56358974	0.05130146
30	1e-02	4.0	0.56358974	0.05130146
31	1e-01	4.0	0.56358974	0.05130146
32	1e+00	4.0	0.47423077	0.08585905
33	5e+00	4.0	0.46160256	0.07545274
34	1e+01	4.0	0.46160256	0.07545274
35	1e+02	4.0	0.46160256	0.07545274

```
bestmod <- tune.out$best.model
summary(bestmod)
```

Call:

```
best.tune(METHOD = svm, train.x = mpg_b ~ . - mpg, data = Auto, ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100), gamma = c(0.5, 1, 2, 3, 4)), kernel = "radial")
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1
```

Number of Support Vectors: 376

( 187 189 )

Number of Classes: 2

Levels:

0 1

```
ypred <- predict(bestmod, Auto)

table(predict = ypred, truth = Auto$mpg_b)
```

	truth	
predict	0	1
0	195	2
1	1	194

```
tune.out <- tune(svm, mpg_b ~ .-mpg, data = Auto, kernel = "poly",
                 ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100),
                                degree = c(1, 2, 3, 4)))
summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost	degree
5	1

- best performance: 0.08935897

- Detailed performance results:

	cost	degree	error	dispersion
1	1e-03	1	0.58634615	0.04634541
2	1e-02	1	0.58634615	0.04634541
3	1e-01	1	0.26455128	0.13497715
4	1e+00	1	0.10987179	0.03254492
5	5e+00	1	0.08935897	0.03884801
6	1e+01	1	0.09192308	0.04407224

7	1e+02	1	0.09192308	0.03264577
8	1e-03	2	0.58634615	0.04634541
9	1e-02	2	0.58634615	0.04634541
10	1e-01	2	0.58634615	0.04634541
11	1e+00	2	0.58634615	0.04634541
12	5e+00	2	0.58634615	0.04634541
13	1e+01	2	0.57352564	0.05444489
14	1e+02	2	0.31134615	0.07418073
15	1e-03	3	0.58634615	0.04634541
16	1e-02	3	0.58634615	0.04634541
17	1e-01	3	0.58634615	0.04634541
18	1e+00	3	0.58634615	0.04634541
19	5e+00	3	0.58634615	0.04634541
20	1e+01	3	0.58634615	0.04634541
21	1e+02	3	0.40237179	0.12105339
22	1e-03	4	0.58634615	0.04634541
23	1e-02	4	0.58634615	0.04634541
24	1e-01	4	0.58634615	0.04634541
25	1e+00	4	0.58634615	0.04634541
26	5e+00	4	0.58634615	0.04634541
27	1e+01	4	0.58634615	0.04634541
28	1e+02	4	0.58634615	0.04634541

```
bestmod <- tune.out$best.model
summary(bestmod)
```

Call:

```
best.tune(METHOD = svm, train.x = mpg_b ~ . - mpg, data = Auto, ranges = list(cost = c(0.001
  0.01, 0.1, 1, 5, 10, 100), degree = c(1, 2, 3, 4)), kernel = "poly")
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: polynomial
  cost: 5
  degree: 1
  coef.0: 0
```

Number of Support Vectors: 148

( 73 75 )

Number of Classes: 2

Levels:

0 1


```
ypred <- predict(bestmod, Auto)

table(predict = ypred, truth = Auto$mpg_b)
```

	truth	
predict	0	1
0	170	9
1	26	187

```
svm.radial <- svm(mpg_b ~ cylinders + displacement + horsepower + weight + acceleration + year,
                  kernel = "radial", cost = 1, gamma = 0.5)
plot(svm.radial, Auto, mpg ~ horsepower)
```

**SVM classification plot**



#8. #a) Create a training set containing a random sample of 800 observations

```
set.seed(1)
indexes <- sample(1:nrow(OJ), 800)
train <- OJ[indexes, ]
test <- OJ[-indexes, ]
```

```
svmfit <- svm(Purchase ~ ., data = train, kernel = "linear",
              cost = 0.01, scale = FALSE)

summary(svmfit)
```

Call:

```
svm(formula = Purchase ~ ., data = train, kernel = "linear", cost = 0.01,
     scale = FALSE)
```

Parameters:

```
SVM-Type:  C-classification
SVM-Kernel: linear
cost:      0.01
```

Number of Support Vectors: 615

```
( 309 306 )
```

Number of Classes: 2

Levels:

```
CH MM
```

```
ypred <- predict(svmfit, train)
table(predict = ypred, truth = train$Purchase)
```

```
      truth
predict CH  MM
CH    420 105
MM     65 210
```

```
 #(65+105)/800 =0.2125
```

```
ypred <- predict(svmfit, test)
table(predict = ypred, truth = test$Purchase)
```

	truth	
predict	CH	MM
CH	148	43
MM	20	59

$\#(20+42)/270 = 0.22962963$

#d) Use the `tune()` function to select an optimal cost. Consider values in the range 0.01 to 10.

```
set.seed(1)
tune.out <- tune(svm, Purchase ~ ., data = train, kernel = "linear",
                 ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
cost
0.1
```

- best performance: 0.1725

- Detailed performance results:

	cost	error	dispersion
1	1e-03	0.31250	0.04124790
2	1e-02	0.17625	0.02853482
3	1e-01	0.17250	0.03162278
4	1e+00	0.17500	0.02946278
5	5e+00	0.17250	0.03162278
6	1e+01	0.17375	0.03197764
7	1e+02	0.17500	0.03486083

```
bestmod <- tune.out$best.model
summary(bestmod)
```

Call:

```
best.tune(METHOD = svm, train.x = Purchase ~ ., data = train, ranges = list(cost = c(0.001,
0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 0.1
```

Number of Support Vectors: 342

```
( 171 171 )
```

Number of Classes: 2

Levels:

```
CH MM
```

```
ypred <- predict(bestmod, train)
table(predict = ypred, truth = train$Purchase)
```

```
      truth
predict CH  MM
CH      422  69
MM      63  246
```

$\#(69+63)/800 = 0.165$

```
ypred <- predict(bestmod, test)
table(predict = ypred, truth = test$Purchase)
```

```
      truth
predict CH  MM
CH      155  31
MM      13  71
```

$\#(13+31)/270 = 0.162962963$

#f) Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for gamma

```

set.seed(1)
tune.out <- tune(svm, Purchase ~ ., data = train,
                 kernel = "radial",
                 ranges = list(cost = c(0.1, 1, 10, 100, 1000),
                               gamma = c(0.5, 1, 2, 3, 4)
                              ) )
summary(tune.out)

```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```

cost gamma
10    0.5

```

- best performance: 0.2125

- Detailed performance results:

	cost	gamma	error	dispersion
1	1e-01	0.5	0.28250	0.05502525
2	1e+00	0.5	0.21375	0.03701070
3	1e+01	0.5	0.21250	0.03632416
4	1e+02	0.5	0.23875	0.04016027
5	1e+03	0.5	0.23875	0.06248611
6	1e-01	1.0	0.34500	0.04937104
7	1e+00	1.0	0.22625	0.04466309
8	1e+01	1.0	0.23000	0.04684490
9	1e+02	1.0	0.24375	0.04973890
10	1e+03	1.0	0.24250	0.05658082
11	1e-01	2.0	0.38625	0.04348132
12	1e+00	2.0	0.22750	0.04281744
13	1e+01	2.0	0.24000	0.04158325
14	1e+02	2.0	0.25875	0.05205833
15	1e+03	2.0	0.26375	0.04910660
16	1e-01	3.0	0.39375	0.04007372
17	1e+00	3.0	0.22625	0.03304563
18	1e+01	3.0	0.25375	0.03335936
19	1e+02	3.0	0.26125	0.03793727
20	1e+03	3.0	0.26375	0.03557562
21	1e-01	4.0	0.39375	0.04007372



```

22 1e+00    4.0 0.22750 0.03322900
23 1e+01    4.0 0.25500 0.03496029
24 1e+02    4.0 0.26250 0.03280837
25 1e+03    4.0 0.26750 0.03073181

```

```

table(
  true = train$Purchase,
  pred = predict(
    tune.out$best.model, newdata = train
  ) )

```

```

      pred
true CH  MM
CH  449   36
MM   49  266

```

$\#(49+36)/800 = 0.10625$

```

table(
  true = test$Purchase,
  pred = predict(
    tune.out$best.model, newdata = test
  ) )

```

```

      pred
true CH  MM
CH  152   16
MM   37   65

```

$\#(37+16)/270 = 0.196296296$

$\#g)$

```

set.seed(1)
tune.out <- tune(svm, Purchase ~ ., data = train,
  kernel = "poly",
  ranges = list(cost = c(0.1, 1, 10, 100, 1000),
    degree = c(0.5, 1, 2, 3, 4)
  ) )
summary(tune.out)

```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost degree  
10 1

- best performance: 0.17125

- Detailed performance results:

	cost	degree	error	dispersion
1	1e-01	0.5	0.39375	0.04007372
2	1e+00	0.5	0.39375	0.04007372
3	1e+01	0.5	0.39375	0.04007372
4	1e+02	0.5	0.39375	0.04007372
5	1e+03	0.5	0.39375	0.04007372
6	1e-01	1.0	0.18000	0.02776389
7	1e+00	1.0	0.17625	0.02853482
8	1e+01	1.0	0.17125	0.02703521
9	1e+02	1.0	0.17375	0.03304563
10	1e+03	1.0	0.17500	0.03486083
11	1e-01	2.0	0.32125	0.05001736
12	1e+00	2.0	0.20250	0.04116363
13	1e+01	2.0	0.18125	0.02779513
14	1e+02	2.0	0.18250	0.02513851
15	1e+03	2.0	0.19125	0.02503470
16	1e-01	3.0	0.28750	0.05068969
17	1e+00	3.0	0.18500	0.02415229
18	1e+01	3.0	0.19500	0.03184162
19	1e+02	3.0	0.22000	0.04609772
20	1e+03	3.0	0.23625	0.04656611
21	1e-01	4.0	0.31875	0.04903584
22	1e+00	4.0	0.23000	0.03016160
23	1e+01	4.0	0.20375	0.02949223
24	1e+02	4.0	0.21375	0.02729087
25	1e+03	4.0	0.23000	0.03016160

```
table(  
  true = train$Purchase,  
  pred = predict(  
    svm,    newdata = test[,1:10],  
    type = "class")
```

```
tune.out$best.model, newdata = train
) )
```

```
      pred
true  CH  MM
CH 424  61
MM  71 244
```

$\#(71+61)/800=0.165$

```
table(
  true = test$Purchase,
  pred = predict(
    tune.out$best.model, newdata = test
  ) )
```

```
      pred
true  CH  MM
CH 155  13
MM  29  73
```

$\#(29+13)/270 = 0.155555556$