

Untitled

```
library(ISLR2)
Gitters <- na.omit(Hitters)
n <- nrow(Gitters)
set.seed(13)
ntest <- trunc(n / 3)
testid <- sample(1:n, ntest)
library(glmnet)
```

Loading required package: Matrix

Loaded glmnet 4.1-8

```
library(torch)
```

Warning: package 'torch' was built under R version 4.3.3

```
library(luz) # high-level interface for torch
library(torchvision) # for datasets and image transformation
```

Warning: package 'torchvision' was built under R version 4.3.3

```
library(torchdatasets) # for datasets we are going to use
```

Warning: package 'torchdatasets' was built under R version 4.3.3

```
library(zeallot)
torch_manual_seed(13)
library(ggplot2)
```

#####labs #10.9.1 A Single Layer Network on the Hitters Data

```
lfit <- lm(Salary ~ ., data = Gitters[-testid, ])
lpred <- predict(lfit, Gitters[testid, ])
with(Gitters[testid, ], mean(abs(lpred - Salary)))
```

[1] 254.6687

fit the lasso

```
x <- scale(model.matrix(Salary ~ . - 1, data = Gitters))
y <- Gitters$Salary

cvfit <- cv.glmnet(x[-testid, ], y[-testid],
  type.measure = "mae")
cpred <- predict(cvfit, x[testid, ], s = "lambda.min")
mean(abs(y[testid] - cpred))
```

[1] 253.1761

#to fit the neural network, we first set up a model structure that describes the network.

algorithm tracks the mean absolute error on the training data, and on validation data if it is supplied.

```
modnn <- nn_module(
  initialize = function(input_size) {
    self$hidden <- nn_linear(input_size, 50)
    self$activation <- nn_relu()
    self$dropout <- nn_dropout(0.2)
    self$output <- nn_linear(50, 1)
  },
```

```

forward = function(x) {
  x %>%
    self$hidden() %>%
    self$activation() %>%
    self$dropout() %>%
    self$output()
}
)


x <- model.matrix(Salary ~ . - 1, data = Gitters) %>% scale()

modnn <- modnn %>%
  setup(
    loss = nn_mse_loss(),
    optimizer = optim_rmsprop,
    metrics = list(luz_metric_mae())
  ) %>%
  set_hparams(input_size = ncol(x))

fitted <- modnn %>%
  fit(
    data = list(x[-testid, ], matrix(y[-testid], ncol = 1)),
    valid_data = list(x[testid, ], matrix(y[testid], ncol = 1)),
    epochs = 50
  )

plot(fitted)

```



```
npred <- predict(fitted, x[testid, ])
npred_array <- as_array(npred)
mae <- mean(abs(y[testid] - npred_array))
print(mae)
```

```
[1] 263.2895
```

```
,
```

#10.9.2 A Multilayer Network on the MNIST Digit Data

```
library(ISLR2)
train_ds <- mnist_dataset(root = ".", train = TRUE, download = TRUE)
test_ds <- mnist_dataset(root = ".", train = FALSE, download = TRUE)

str(train_ds[1])
```

```
List of 2
```

```
$ x: int [1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 ...
```

```
$ y: int 6
```

```
str(test_ds[2])
```

List of 2

```
$ x: int [1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 ...
```

```
$ y: int 3
```

```
length(train_ds)
```

```
[1] 60000
```

```
length(test_ds)
```

```
[1] 10000
```

```
transform <- function(x) {  
  x %>%  
    torch_tensor() %>%  
    torch_flatten() %>%  
    torch_div(255)  
}  
train_ds <- mnist_dataset(  
  root = ".",  
  train = TRUE,  
  download = TRUE,  
  transform = transform  
)  
test_ds <- mnist_dataset(  
  root = ".",  
  train = FALSE,  
  download = TRUE,  
  transform = transform  
)  
  
modelnn <- nn_module(  
  initialize = function() {  
    self$linear1 <- nn_linear(in_features = 28*28, out_features = 256)  
    self$linear2 <- nn_linear(in_features = 256, out_features = 128)
```

```

self$linear3 <- nn_linear(in_features = 128, out_features = 10)

self$drop1 <- nn_dropout(p = 0.4)
self$drop2 <- nn_dropout(p = 0.3)

self$activation <- nn_relu()
},
forward = function(x) {
  x %>%

    self$linear1() %>%
    self$activation() %>%
    self$drop1() %>%

    self$linear2() %>%
    self$activation() %>%
    self$drop2() %>%

    self$linear3()
  }
)

print(modelnn())

```

An `nn_module` containing 235,146 parameters.

```

-- Modules -----
* linear1: <nn_linear> #200,960 parameters
* linear2: <nn_linear> #32,896 parameters
* linear3: <nn_linear> #1,290 parameters
* drop1: <nn_dropout> #0 parameters
* drop2: <nn_dropout> #0 parameters
* activation: <nn_relu> #0 parameters

modelnn <- modelnn %>%
  setup(
    loss = nn_cross_entropy_loss(),
    optimizer = optim_rmsprop,
    metrics = list(luz_metric_accuracy())
  )

```

```

system.time(
  fitted <- modelnn %>%
    fit(
      data = train_ds,
      epochs = 5,
      valid_data = 0.2,
      dataloader_options = list(batch_size = 256),
      verbose = FALSE
    )
)


```

```

user  system elapsed
34.556  1.980  36.751

```

```
plot(fitted)
```



```

accuracy <- function(pred, truth) {
  mean(pred == truth) }

# gets the true classes from all observations in test_ds.

```

```

truth <- sapply(seq_along(test_ds), function(x) test_ds[x][[2]])

fitted %>%
  predict(test_ds) %>%
  torch_argmax(dim = 2) %>% # the predicted class is the one with higher 'logit'.
  as_array() %>% # we convert to an R object
  accuracy(truth)

```

```
[1] 0.9532
```

```

modellr <- nn_module(
  initialize = function() {
    self$linear <- nn_linear(784, 10)
  },
  forward = function(x) {
    self$linear(x)
  }
)
print(modellr())

```

An `nn_module` containing 7,850 parameters.

```

-- Modules -----
* linear: <nn_linear> #7,850 parameters

```

```

fit_modellr <- modellr %>%
  setup(
    loss = nn_cross_entropy_loss(),
    optimizer = optim_rmsprop,
    metrics = list(luz_metric_accuracy())
  ) %>%
  fit(
    data = train_ds,
    epochs = 5,
    valid_data = 0.2,
    dataloader_options = list(batch_size = 128)
  )

fit_modellr %>%

```



```

predict(test_ds) %>%
  torch_argmax(dim = 2) %>% # the predicted class is the one with higher 'logit'.
  as_array() %>% # we convert to an R object
  accuracy(truth)

```

```
[1] 0.918
```

```

# alternatively one can use the `evaluate` function to get the results
# on the test_ds
evaluate(fit_modelldr, test_ds)

```

```
A `luz_module_evaluation`
```

```

-- Results -----
loss: 0.3
acc: 0.918

```

#10.9.3 Convolutional Neural Networks

```

transform <- function(x) {
  transform_to_tensor(x)
}

train_ds <- cifar100_dataset(
  root = "./",
  train = TRUE,
  download = TRUE,
  transform = transform
)

test_ds <- cifar100_dataset(
  root = "./",
  train = FALSE,
  transform = transform
)

str(train_ds[1])

```

```
List of 2
```

```


$ x:Float [1:3, 1:32, 1:32]
$ y: int 20

```

```
length(train_ds)
```

```
[1] 50000
```

```
par(mar = c(0, 0, 0, 0), mfrow = c(5, 5))
index <- sample(seq(50000), 25)
for (i in index) plot(as.raster(as.array(train_ds[i][[1]]$permute(c(2,3,1)))))
```



```
conv_block <- nn_module(  
  initialize = function(in_channels, out_channels) {  
    self$conv <- nn_conv2d(  
      in_channels = in_channels,  
      out_channels = out_channels,  
      kernel_size = c(3,3),  
      padding = "same"  
    )  
    self$relu <- nn_relu()  
    self$pool <- nn_max_pool2d(kernel_size = c(2,2))  
  },  
  forward = function(x) {
```

```

        x %>%
        self$conv() %>%
        self$relu() %>%
        self$pool()
    }
)

model <- nn_module(
  initialize = function() {
    self$conv <- nn_sequential(
      conv_block(3, 32),
      conv_block(32, 64),
      conv_block(64, 128),
      conv_block(128, 256)
    )
    self$output <- nn_sequential(
      nn_dropout(0.5),
      nn_linear(2*2*256, 512),
      nn_relu(),
      nn_linear(512, 100)
    )
  },
  forward = function(x) {
    x %>%
    self$conv() %>%
    torch_flatten(start_dim = 2) %>%
    self$output()
  }
)
model()

```

An `nn_module` containing 964,516 parameters.

```

-- Modules -----
* conv: <nn_sequential> #388,416 parameters
* output: <nn_sequential> #576,100 parameters

```

```

fitted <- model %>%
  setup(
    loss = nn_cross_entropy_loss(),
    optimizer = optim_rmsprop,

```

```

        metrics = list(luz_metric_accuracy())
    ) %>%
    set_opt_hparams(lr = 0.001) %>%
    fit(
        train_ds,
        epochs = 10, #30,
        valid_data = 0.2,
        dataloader_options = list(batch_size = 128)
    )

print(fitted)

```

A `luz_module_fitted`

```

-- Time -----
* Total time: 3m 17.9s
* Avg time per training epoch: 17.1s

-- Results -----
Metrics observed in the last epoch.

i Training:
loss: 2.351
acc: 0.3824

-- Model -----
An `nn_module` containing 964,516 parameters.

-- Modules -----
* conv: <nn_sequential> #388,416 parameters
* output: <nn_sequential> #576,100 parameters

```

```

evaluate(fitted, test_ds)

```

A `luz_module_evaluation`

```

-- Results -----
loss: 2.4013
acc: 0.3816

```

#10.9.4 Using Pretrained CNN Models

```

img_dir <- "book_images"
image_names <- list.files(img_dir)
num_images <- length(image_names)
x <- torch_empty(num_images, 3, 224, 224)
for (i in 1:num_images) {
  img_path <- file.path(img_dir, image_names[i])
  img <- img_path %>%
    base_loader() %>%
    transform_to_tensor() %>%
    transform_resize(c(224, 224)) %>%
    # normalize with imagenet mean and stds.
    transform_normalize(
      mean = c(0.485, 0.456, 0.406),
      std = c(0.229, 0.224, 0.225)
    )
  x[i,, ] <- img
}

model <- torchvision::model_resnet18(pretrained = TRUE)
model$eval() # put the model in evaluation mode
preds <- model(x)

mapping <- jsonlite::read_json("https://s3.amazonaws.com/deep-learning-models/image-models/
  sapply(function(x) x[[2]])

top3 <- torch_topk(preds, dim = 2, k = 3)

top3_prob <- top3[[1]] %>%
  nnf_softmax(dim = 2) %>%
  torch_unbind() %>%
  lapply(as.numeric)

top3_class <- top3[[2]] %>%
  torch_unbind() %>%
  lapply(function(x) mapping[as.integer(x)])

result <- purrr::map2(top3_prob, top3_class, function(pr, cl) {
  names(pr) <- cl
  pr
})
names(result) <- image_names

```

```
print(result)
```

```
$Cape_Weaver.jpg
hummingbird    lorikeet    bee_eater
0.3633293     0.3577291    0.2789416
```

```
$Hawk_cropped.jpg
kite           jay         magpie
0.6157786     0.2311880    0.1530334
```

```
$Hawk_Fountain.jpg
eel            agama    common_newt
0.5391121     0.2527187    0.2081692
```

```
$Lhasa_Apso.jpg
Lhasa Tibetan_terrier    Shih-Tzu
0.79760498    0.12012957    0.08226541
```

```
$Sleeping_Cat.jpg
Saint_Bernard    guinea_pig    Bernese_mountain_dog
0.3946666        0.3426994        0.2626340
```

#10.9.5 IMDb Document Classification

```
max_features <- 10000
imdb_train <- imdb_dataset(
  root = ".",
  download = TRUE,
  num_words = max_features
)
imdb_test <- imdb_dataset(
  root = ".",
  download = TRUE,
  num_words = max_features
)

imdb_train[1]$x[1:12]
```

```
[1] 2 261 297 14 20 23 4 6253 1307 13 70 65
```

```

word_index <- imdb_train$vocabulary
decode_review <- function(text, word_index) {
  word <- names(word_index)
  idx <- unlist(word_index, use.names = FALSE)
  word <- c("<PAD>", "<START>", "<UNK>", word)
  words <- word[text]
  paste(words, collapse = " ")
}
decode_review(imdb_train[1]$x[1:12], word_index)

```

[1] "<START> having watched this movie on the scifi channel i can only"

```

library(Matrix)
one_hot <- function(sequences, dimension) {
  seqlen <- sapply(sequences, length)
  n <- length(seqlen)
  rowind <- rep(1:n, seqlen)
  colind <- unlist(sequences)
  sparseMatrix(i = rowind, j = colind,
    dims = c(n, dimension))
}

# collect all values into a list
train <- seq_along(imdb_train) %>%
  lapply(function(i) imdb_train[i]) %>%
  purrr::transpose()
test <- seq_along(imdb_test) %>%
  lapply(function(i) imdb_test[i]) %>%
  purrr::transpose()

# num_words + padding + start + oov token = 10000 + 3
x_train_1h <- one_hot(train$x, 10000 + 3)
x_test_1h <- one_hot(test$x, 10000 + 3)
dim(x_train_1h)

```

[1] 25000 10003

```

nnzero(x_train_1h) / (25000 * (10000 + 3))

```


```
[1] 0.01316756
```

```
set.seed(3)
ival <- sample(seq(along = train$y), 2000)
itrain <- seq_along(train$y)[-ival]

library(glmnet)
y_train <- unlist(train$y)

fitlm <- glmnet(x_train_1h[itrain, ], unlist(y_train[itrain]),
  family = "binomial", standardize = FALSE)
classlmv <- predict(fitlm, x_train_1h[ival, ]) > 0
acclmv <- apply(classlmv, 2, accuracy, unlist(y_train[ival]) > 0)

par(mar = c(4, 4, 4, 4), mfrow = c(1, 1))
plot(-log(fitlm$lambda), acclmv)
```



```
model <- nn_module(
  initialize = function(input_size = 10000 + 3) {
    self$dense1 <- nn_linear(input_size, 16)
    self$relu <- nn_relu()
```




```

        self$dense2 <- nn_linear(16, 16)
        self$output <- nn_linear(16, 1)
    },
    forward = function(x) {
        x %>%
            self$dense1() %>%
            self$relu() %>%
            self$dense2() %>%
            self$relu() %>%
            self$output() %>%
            torch_flatten(start_dim = 1)
    }
)
model <- model %>%
    setup(
        loss = nn_bce_with_logits_loss(),
        optimizer = optim_rmsprop,
        metrics = list(luz_metric_binary_accuracy_with_logits())
    ) %>%
    set_opt_hparams(lr = 0.001)

fitted <- model %>%
    fit(
        # we transform the training and validation data into torch tensors
        list(
            torch_tensor(as.matrix(x_train_1h[itrain,]), dtype = torch_float()),
            torch_tensor(unlist(train$y[itrain]))
        ),
        valid_data = list(
            torch_tensor(as.matrix(x_train_1h[ival, ]), dtype = torch_float()),
            torch_tensor(unlist(train$y[ival]))
        ),
        dataloader_options = list(batch_size = 512),
        epochs = 10
    )

plot(fitted)

```



```
fitted <- model %>%
  fit(
    list(
      torch_tensor(as.matrix(x_train_1h[itrain,]), dtype = torch_float()),
      torch_tensor(unlist(train$y[itrain]))
    ),
    valid_data = list(
      torch_tensor(as.matrix(x_test_1h), dtype = torch_float()),
      torch_tensor(unlist(test$y))
    ),
    dataloader_options = list(batch_size = 512),
    epochs = 10
  )
```

#10.9.6 Recurrent Neural Networks ## Sequential Models for Document Classification

```
wc <- sapply(seq_along(imdb_train), function(i) length(imdb_train[i]$x))
median(wc)
```

[1] 178

```
sum(wc <= 500) / length(wc)
```

```
[1] 0.916
```

```
maxlen <- 500
num_words <- 10000
imdb_train <- imdb_dataset(root = ".", split = "train", num_words = num_words,
                           maxlen = maxlen)
imdb_test <- imdb_dataset(root = ".", split = "test", num_words = num_words,
                          maxlen = maxlen)

vocab <- c(rep(NA, imdb_train$index_from - 1), imdb_train$get_vocabulary())
tail(names(vocab)[imdb_train[1]$x])
```

```
[1] "compensate" "you"          "the"          "rental"       ""
[6] "d"
```

```
model <- nn_module(
  initialize = function() {
    self$embedding <- nn_embedding(10000 + 3, 32)
    self$lstm <- nn_lstm(input_size = 32, hidden_size = 32, batch_first = TRUE)
    self$dense <- nn_linear(32, 1)
  },
  forward = function(x) {
    c(output, c(hn, cn)) %<-% (x %>%
      self$embedding() %>%
      self$lstm())
    output[,-1,] %>% # get the last output
      self$dense() %>%
      torch_flatten(start_dim = 1)
  }
)

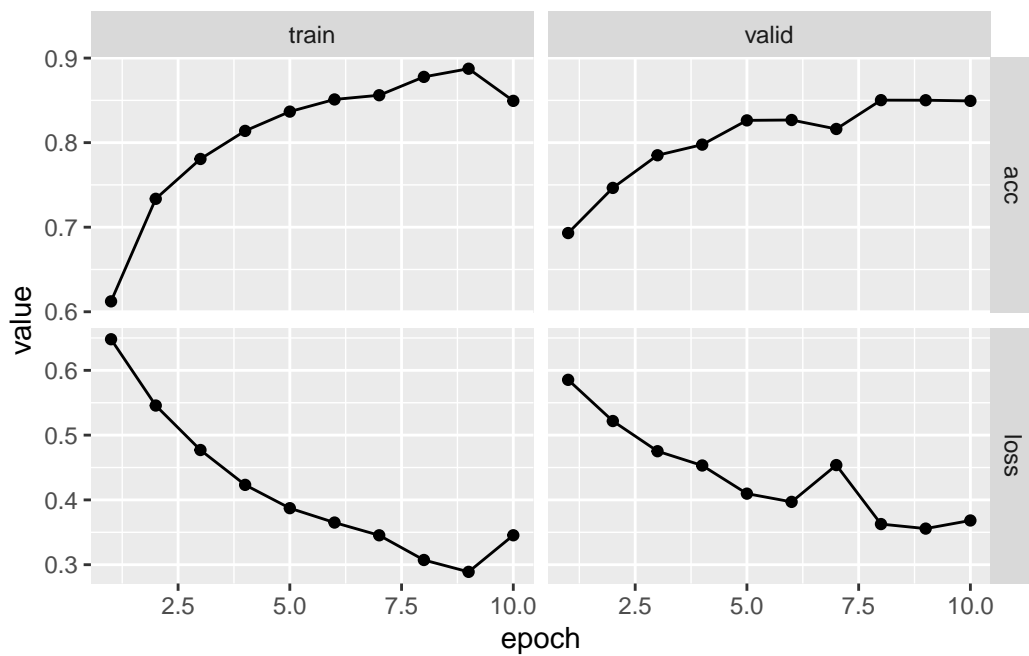
model <- model %>%
  setup(
    loss = nn_bce_with_logits_loss(),
    optimizer = optim_rmsprop,
    metrics = list(luz_metric_binary_accuracy_with_logits())
  ) %>%
```

```

set_opt_hparams(lr = 0.001)

fitted <- model %>% fit(
  imdb_train,
  epochs = 10,
  dataloader_options = list(batch_size = 128),
  valid_data = imdb_test
)
plot(fitted)

```



```

predy <- torch_sigmoid(predict(fitted, imdb_test)) > 0.5
evaluate(fitted, imdb_test, dataloader_options = list(batch_size = 512))

```

A `luz_module_evaluation`

-- Results -----

loss: 0.3685

acc: 0.8494

Time Series Prediction

```
library(ISLR2)
xdata <- data.matrix(
  NYSE[, c("DJ_return", "log_volume", "log_volatility")]
)
istrain <- NYSE[, "train"]
xdata <- scale(xdata)

lagm <- function(x, k = 1) {
  n <- nrow(x)
  pad <- matrix(NA, k, ncol(x))
  rbind(pad, x[1:(n - k), ])
}

arframe <- data.frame(log_volume = xdata[, "log_volume"],
  L1 = lagm(xdata, 1), L2 = lagm(xdata, 2),
  L3 = lagm(xdata, 3), L4 = lagm(xdata, 4),
  L5 = lagm(xdata, 5)
)

arframe <- arframe[-(1:5), ]
istrain <- istrain[-(1:5)]

arfit <- lm(log_volume ~ ., data = arframe[istrain, ])
arpred <- predict(arfit, arframe[!istrain, ])
V0 <- var(arframe[!istrain, "log_volume"])
1 - mean((arpred - arframe[!istrain, "log_volume"])^2) / V0
```

[1] 0.413223

```
arframed <-
  data.frame(day = NYSE[-(1:5), "day_of_week"], arframe)
arfitd <- lm(log_volume ~ ., data = arframed[istrain, ])
arpredd <- predict(arfitd, arframed[!istrain, ])
1 - mean((arpredd - arframe[!istrain, "log_volume"])^2) / V0
```

[1] 0.4598616

```

n <- nrow(arframe)
xrnn <- data.matrix(arframe[, -1])
xrnn <- array(xrnn, c(n, 3, 5))
xrnn <- xrnn[, , 5:1]
xrnn <- aperm(xrnn, c(1, 3, 2))
dim(xrnn)

```

```
[1] 6046      5      3
```

```

model <- nn_module(
  initialize = function() {
    self$rnn <- nn_rnn(3, 12, batch_first = TRUE)
    self$dense <- nn_linear(12, 1)
    self$dropout <- nn_dropout(0.2)
  },
  forward = function(x) {
    c(output, ...) %<-% (x %>%
      self$rnn())
    output[, -1, ] %>%
      self$dropout() %>%
      self$dense() %>%
      torch_flatten(start_dim = 1)
  }
)

model <- model %>%
  setup(
    optimizer = optim_rmsprop,
    loss = nn_mse_loss()
  ) %>%
  set_opt_hparams(lr = 0.001)

fitted <- model %>% fit(
  list(xrnn[istrain, , ], arframe[istrain, "log_volume"]),
  epochs = 75, #epochs = 200,
  dataloader_options = list(batch_size = 64),
  valid_data =
    list(xrnn[!istrain, , ], arframe[!istrain, "log_volume"])
)
kpred <- as.numeric(predict(fitted, xrnn[!istrain, , ]))

```

```
1 - mean((kpred - arframe[!istrain, "log_volume"])^2) / V0
```

```
[1] 0.4131635
```

```
model <- nn_module(
  initialize = function() {
    self$dense <- nn_linear(15, 1)
  },
  forward = function(x) {
    x %>%
      torch_flatten(start_dim = 2) %>%
      self$dense()
  }
)
```

```
x <- model.matrix(log_volume ~ . - 1, data = arframed)
colnames(x)
```

```
[1] "dayfri"          "daymon"          "daythur"
[4] "daytues"         "daywed"          "L1.DJ_return"
[7] "L1.log_volume"   "L1.log_volatility" "L2.DJ_return"
[10] "L2.log_volume"   "L2.log_volatility" "L3.DJ_return"
[13] "L3.log_volume"   "L3.log_volatility" "L4.DJ_return"
[16] "L4.log_volume"   "L4.log_volatility" "L5.DJ_return"
[19] "L5.log_volume"   "L5.log_volatility"
```


```
arnnd <- nn_module(
  initialize = function() {
    self$dense <- nn_linear(15, 32)
    self$dropout <- nn_dropout(0.5)
    self$activation <- nn_relu()
    self$output <- nn_linear(32, 1)
  },
  forward = function(x) {
    x %>%
      torch_flatten(start_dim = 2) %>%
      self$dense() %>%
      self$activation() %>%
      self$output()
  }
)
```

```

        self$dropout() %>%
        self$output() %>%
        torch_flatten(start_dim = 1)
    }
)
arnnd <- arnnd %>%
  setup(
    optimizer = optim_rmsprop,
    loss = nn_mse_loss()
  ) %>%
  set_opt_hparams(lr = 0.001)

fitted <- arnnd %>% fit(
  list(xrnn[istrain,, ], arframe[istrain, "log_volume"]),
  epochs = 30, #epochs = 200,
  dataloader_options = list(batch_size = 64),
  valid_data =
    list(xrnn[!istrain,, ], arframe[!istrain, "log_volume"])
)
plot(fitted)

```




```

npred <- as.numeric(predict(fitted, xrn[!istrain, ]))
1 - mean((arframe[!istrain, "log_volume"] - npred)^2) / V0

```

[1] 0.428227

#6.

```


R_beta <- function(beta) {
  sin(beta) + beta / 10
}

b <- seq(-6, 6, length.out = 400)

R_b <- R_beta(b)

plot(b, R_b, type = "l", col = "blue", lwd = 2,
     xlab = expression(beta), ylab = expression(R(beta)))

```



```

dR_beta <- function(beta) {
  cos(beta) + 0.1
}

```

```

}


beta_0 <- 2.3
learning_rate <- 0.1
num_iterations <- 100

beta_values <- numeric(num_iterations)
beta_values[1] <- beta_0

for (t in 2:num_iterations) {
  beta_values[t] <- beta_values[t-1] - learning_rate * dR_beta(beta_values[t-1])
}

plot(b, R_b, type = "l", col = "blue", lwd = 2, xlab = expression(beta), ylab = expression(R(beta)))
points(beta_values, R_beta(beta_values), col = "red", pch = 19)

```



```

cat("Final value after gradient descent:", beta_values[num_iterations], "\n")

```

Final value after gradient descent: 4.612013


```

beta_1 <- 1.4
beta_values <- numeric(num_iterations)
beta_values[1] <- beta_1

for (t in 2:num_iterations) {
  beta_values[t] <- beta_values[t-1] - learning_rate * dR_beta(beta_values[t-1])
}

plot(b, R_b, type = "l", col = "blue", lwd = 2, xlab = expression(beta), ylab = expression(R(beta)))
points(beta_values, R_beta(beta_values), col = "green", pch = 19)

```



```

cat("Final value after gradient descent:", beta_values[num_iterations], "\n")

```

Final value after gradient descent: -1.670396

#7. Fit a neural network to the Default data.

```

Default$default <- ifelse(Default$default == "Yes", 1, 0)

n <- nrow(Default)

```

```

set.seed(13)
ntest <- trunc(n / 5)
testid <- sample(1:n, ntest)

modnn <- nn_module(
  initialize = function(input_size) {
    self$hidden <- nn_linear(input_size, 10)
    self$activation <- nn_relu()
    self$dropout <- nn_dropout(0.3)
    self$output <- nn_linear(10, 1)
  },
  forward = function(x) {
    x %>%
      self$hidden() %>%
      self$activation() %>%
      self$dropout() %>%
      self$output()
  }
)

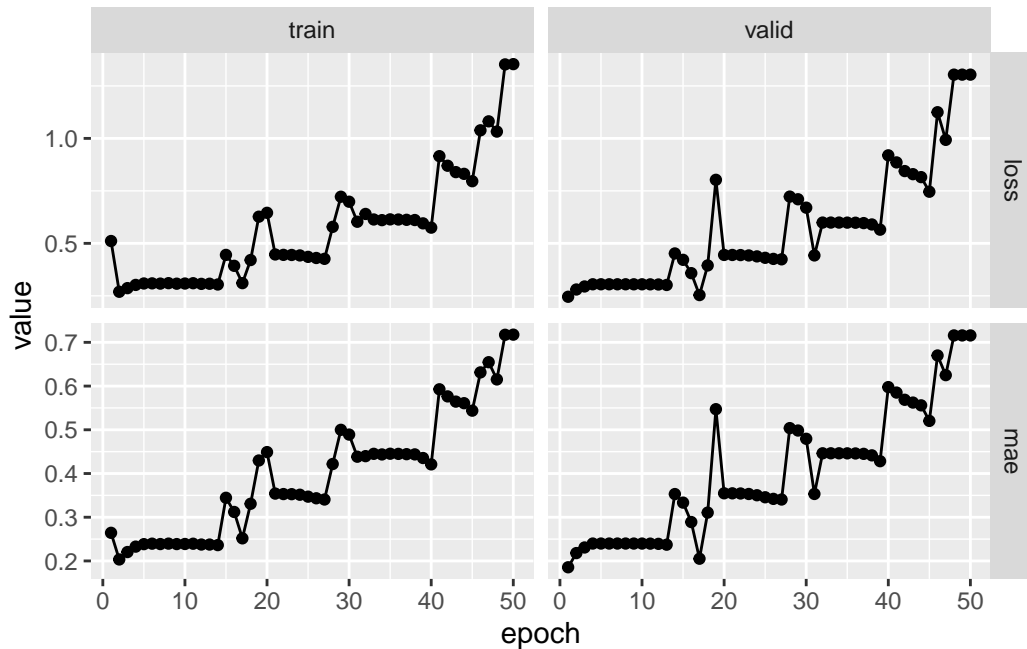
x <- model.matrix(default ~ . - 1, data = Default) %>% scale()
y <- Default$default

modnn <- modnn %>%
  setup(
    loss = nn_bce_loss(),
    optimizer = optim_rmsprop,
    metrics = list(luz_metric_mae())
  ) %>%
  set_hparams(input_size = ncol(x))

fitted <- modnn %>%
  fit(
    data = list(x[-testid, ], matrix(y[-testid], ncol = 1)),
    valid_data = list(x[testid, ], matrix(y[testid], ncol = 1)),
    epochs = 50
  )

plot(fitted)

```



```

npred <- predict(fitted, x[testid, ])
npred_array <- as_array(npred)
predictions <- ifelse(npred_array > 0.5, 1, 0)
accuracy <- mean(predictions == y[testid])

print(paste("Accuracy:", accuracy))

```

```
[1] "Accuracy: 0.0355"
```

#8. From your collection of personal photographs, pick 10 images of animals

```

img_dir <- "animal_images"
image_names <- list.files(img_dir)
num_images <- length(image_names)
x <- torch_empty(num_images, 3, 224, 224)
for (i in 1:num_images) {
  img_path <- file.path(img_dir, image_names[i])
  img <- img_path %>%
    base_loader() %>%
    transform_to_tensor() %>%
    transform_resize(c(224, 224)) %>%

```

```

    # normalize with imagenet mean and stds.
    transform_normalize(
      mean = c(0.485, 0.456, 0.406),
      std = c(0.229, 0.224, 0.225)
    )
    x[i,,, ] <- img
  }

model <- torchvision::model_resnet18(pretrained = TRUE)
model$eval() # put the model in evaluation mode
preds <- model(x)

mapping <- jsonlite::read_json("https://s3.amazonaws.com/deep-learning-models/image-models/
  sapply(function(x) x[[2]])

top5 <- torch_topk(preds, dim = 2, k = 5)

top5_prob <- top5[[1]] %>%
  nnf_softmax(dim = 2) %>%
  torch_unbind() %>%
  lapply(as.numeric)

top5_class <- top5[[2]] %>%
  torch_unbind() %>%
  lapply(function(x) mapping[as.integer(x)])

result <- purrr::map2(top5_prob, top5_class, function(pr, cl) {
  names(pr) <- cl
  pr
})
names(result) <- image_names
print(result)

```

\$IMG_4037.JPG

Labrador_retriever	golden_retriever	cocker_spaniel	beagle
0.648640752	0.321777314	0.015168818	0.008276871
kuvasz			
0.006136307			

\$IMG_4191.JPG

carton	guinea_pig	Band_Aid crossword_puzzle
--------	------------	---------------------------

0.2440549	0.2263593	0.1907472	0.1806872
plastic_bag			
0.1581516			

\$IMG_4588.JPG

Shih-Tzu	Lhasa	Maltese_dog	Pekinese	toy_poodle
0.55186415	0.16366476	0.12603915	0.07993906	0.07849292

\$IMG_4591.JPG

otter	marmot	weasel	mink	beaver
0.36287266	0.34405339	0.14921607	0.08073228	0.06312557

\$IMG_4592.JPG

chickadee	puffer	red-backed_sandpiper
0.40517688	0.32740453	0.09740868
quail	sea_slug	
0.08738206	0.08262786	

\$IMG_4593.JPG

brown_bear	bison	teddy	mongoose	chow
0.65225726	0.19173110	0.07481807	0.04515351	0.03604012

\$IMG_4594.JPG

dhole	cheetah	dingo	standard_poodle
0.2736522	0.2573565	0.2295044	0.1486792
miniature_poodle			
0.0908076			

\$IMG_4595.JPG

Maltese_dog	Persian_cat
0.54989856	0.22765583
Shih-Tzu	West_Highland_white_terrier
0.09817947	0.06886514
Old_English_sheepdog	
0.05540101	

\$IMG_4596.JPG

llama	Angora	kuvasz	guinea_pig	wallaby
0.2987915	0.1975946	0.1930207	0.1647692	0.1458240

\$IMG_4597.JPG

Angora	Persian_cat	Siamese_cat	wood_rabbit	hamster
0.71234006	0.16844420	0.06128484	0.02935519	0.02857562

#9. Fit a lag-5 autoregressive model to the NYSE data

```
library(ISLR2)
xdata <- data.matrix(
  NYSE[, c("DJ_return", "log_volume", "log_volatility")]
)
istrain <- NYSE[, "train"]
xdata <- scale(xdata)

lagm <- function(x, k = 1) {
  n <- nrow(x)
  pad <- matrix(NA, k, ncol(x))
  rbind(pad, x[1:(n - k), ])
}

arframe <- data.frame(log_volume = xdata[, "log_volume"],
  L1 = lagm(xdata, 1), L2 = lagm(xdata, 2),
  L3 = lagm(xdata, 3), L4 = lagm(xdata, 4),
  L5 = lagm(xdata, 5)
)

arframe <- arframe[-(1:5), ]
istrain <- istrain[-(1:5)]
V0 <- var(arframe[!istrain, "log_volume"])
```

```
NYSE$month <- format(as.Date(NYSE$date), "%m")
NYSE$month <- factor(NYSE$month, levels = sprintf("%02d", 1:12))
arframed <- data.frame(day = NYSE[-(1:5), "day_of_week"], month = NYSE[-(1:5), "month"], a

arfitd <- lm(log_volume ~ ., data = arframed[istrain, ])
arpredd <- predict(arfitd, arframed[!istrain, ])
1 - mean((arpredd - arframe[!istrain, "log_volume"])^2) / V0
```

[1] 0.4629872

#0.4629872

#10.

```
n <- nrow(arframe)
xrnn <- data.matrix(arframe[, -1])
```



```

xrnn <- array(xrnn, c(n, 3, 5))
xrnn <- xrnn[, , 5:1]
xrnn <- aperm(xrnn, c(1, 3, 2))
dim(xrnn)

```

```
[1] 6046      5      3
```

```

model <- nn_module(
  initialize = function() {
    self$rnn <- nn_rnn(3, 12, batch_first = TRUE)
    self$dense <- nn_linear(12, 1)
    self$dropout <- nn_dropout(0.2)
  },
  forward = function(x) {
    c(output, ...) %<-% (x %>%
      self$rnn())
    output[,-1,] %>%
      self$dropout() %>%
      self$dense() %>%
      torch_flatten(start_dim = 1)
  }
)

model <- model %>%
  setup(
    optimizer = optim_rmsprop,
    loss = nn_mse_loss()
  ) %>%
  set_opt_hparams(lr = 0.001)

fitted <- model %>% fit(
  list(xrnn[istrain,, ], arframe[istrain, "log_volume"]),
  epochs = 200, #epochs = 200,
  dataloader_options = list(batch_size = 64),
  valid_data =
    list(xrnn[!istrain,, ], arframe[!istrain, "log_volume"])
)

kpred <- as.numeric(predict(fitted, xrnn[!istrain,, ]))
1 - mean((kpred - arframe[!istrain, "log_volume"])^2) / V0

```

```
[1] 0.4085854
```

#11.

```
model <- nn_module(  
  initialize = function() {  
    self$dense <- nn_linear(15, 1)  
  },  
  forward = function(x) {  
    x %>%  
      torch_flatten(start_dim = 2) %>%  
      self$dense()  
  }  
)
```

```
x <- model.matrix(log_volume ~ . - 1, data = arframed)  
colnames(x)
```

```
[1] "dayfri"          "daymon"          "daythur"  
[4] "daytues"         "daywed"          "month02"  
[7] "month03"         "month04"         "month05"  
[10] "month06"         "month07"         "month08"  
[13] "month09"         "month10"         "month11"  
[16] "month12"         "L1.DJ_return"    "L1.log_volume"  
[19] "L1.log_volatility" "L2.DJ_return"    "L2.log_volume"  
[22] "L2.log_volatility" "L3.DJ_return"    "L3.log_volume"  
[25] "L3.log_volatility" "L4.DJ_return"    "L4.log_volume"  
[28] "L4.log_volatility" "L5.DJ_return"    "L5.log_volume"  
[31] "L5.log_volatility"
```


```
arnnd <- nn_module(  
  initialize = function() {  
    self$dense <- nn_linear(15, 32)  
    self$dropout <- nn_dropout(0.5)  
    self$activation <- nn_relu()  
    self$output <- nn_linear(32, 1)  
  },  
  forward = function(x) {  
    x %>%  
      torch_flatten(start_dim = 2) %>%  
      self$dense() %>%  
      self$activation() %>%  
      self$output()  
  }  
)
```

```

        self$dropout() %>%
        self$output() %>%
        torch_flatten(start_dim = 1)
    }
)
arnnd <- arnnd %>%
  setup(
    optimizer = optim_rmsprop,
    loss = nn_mse_loss()
  ) %>%
  set_opt_hparams(lr = 0.001)

fitted <- arnnd %>% fit(
  list(xrnn[istrain,, ], arframe[istrain, "log_volume"]),
  epochs = 30,
  dataloader_options = list(batch_size = 64),
  valid_data =
    list(xrnn[!istrain,, ], arframe[!istrain, "log_volume"])
)
plot(fitted)

```



```

npred <- as.numeric(predict(fitted, xrn[!istrain, ]))
1 - mean((arframe[!istrain, "log_volume"] - npred)^2) / V0

```

```
[1] 0.4277066
```

#13.

```

dict_sizes <- c(1000, 3000, 5000, 10000)

accuracy <- function(pred, truth) {
  mean(pred == truth) }

# Loop through each dictionary size
for (max_features in dict_sizes) {

  cat("\n\nTesting with dictionary size:", max_features, "\n\n")

  imdb_train <- imdb_dataset(
    root = ".",
    download = TRUE,
    num_words = max_features
  )

  imdb_test <- imdb_dataset(
    root = ".",
    download = TRUE,
    num_words = max_features
  )

  word_index <- imdb_train$vocabulary

  decode_review <- function(text, word_index) {
    word <- names(word_index)
    idx <- unlist(word_index, use.names = FALSE)
    word <- c("<PAD>", "<START>", "<UNK>", word)
    words <- word[text]
    paste(words, collapse = " ")
  }

  print(decode_review(imdb_train[1]$x[1:12], word_index))

```

```

one_hot <- function(sequences, dimension) {
  seqlen <- sapply(sequences, length)
  n <- length(seqlen)
  rowind <- rep(1:n, seqlen)
  colind <- unlist(sequences)
  sparseMatrix(i = rowind, j = colind, dims = c(n, dimension))
}

train <- seq_along(imdb_train) %>%
  lapply(function(i) imdb_train[i]) %>%
  purrr::transpose()

test <- seq_along(imdb_test) %>%
  lapply(function(i) imdb_test[i]) %>%
  purrr::transpose()

# One-hot encoding (adjust the size according to max_features)
x_train_1h <- one_hot(train$x, max_features + 3) # Padding and special tokens
x_test_1h <- one_hot(test$x, max_features + 3)

cat("Train data dimensions: ", dim(x_train_1h), "\n")
cat("Non-zero elements in train set: ", nnzero(x_train_1h) / (25000 * (max_features + 3))

set.seed(3)
ival <- sample(seq(along = train$y), 2000) # Validation set indices
itrain <- seq_along(train$y)[-ival] # Training set indices

y_train <- unlist(train$y)

fitlm <- glmnet(x_train_1h[itrain, ], unlist(y_train[itrain]),
  family = "binomial", standardize = FALSE)

classlmv <- predict(fitlm, x_train_1h[ival, ]) > 0
acclmv <- apply(classlmv, 2, accuracy, unlist(y_train[ival]) > 0)

model <- nn_module(
  initialize = function(input_size = max_features + 3) {
    self$dense1 <- nn_linear(input_size, 16)
  }

```

```

    self$relu <- nn_relu()
    self$dense2 <- nn_linear(16, 16)
    self$output <- nn_linear(16, 1)
  },
  forward = function(x) {
    x %>%
      self$dense1() %>%
      self$relu() %>%
      self$dense2() %>%
      self$relu() %>%
      self$output() %>%
      torch_flatten(start_dim = 1)
  }
)

model <- model %>%
  setup(
    loss = nn_bce_with_logits_loss(),
    optimizer = optim_rmsprop,
    metrics = list(luz_metric_binary_accuracy_with_logits())
  ) %>%
  set_opt_hparams(lr = 0.001)

fitted <- model %>%
  fit(
    list(
      torch_tensor(as.matrix(x_train_1h[itrain, ]), dtype = torch_float()),
      torch_tensor(unlist(train$y[itrain]))
    ),
    valid_data = list(
      torch_tensor(as.matrix(x_train_1h[ival, ]), dtype = torch_float()),
      torch_tensor(unlist(train$y[ival]))
    ),
    dataloader_options = list(batch_size = 512),
    epochs = 10
  )

fitted <- model %>%
  fit(
    list(

```

```

        torch_tensor(as.matrix(x_train_1h[itrain, ]), dtype = torch_float()),
        torch_tensor(unlist(train$y[itrain]))
    ),
    valid_data = list(
        torch_tensor(as.matrix(x_test_1h), dtype = torch_float()),
        torch_tensor(unlist(test$y))
    ),
    dataloader_options = list(batch_size = 512),
    epochs = 10
)

cat("\nFinished testing with dictionary size:", max_features, "\n\n")
}

```

Testing with dictionary size: 1000

```

[1] "<START> this does give away some of the plot by the way"
Train data dimensions: 25000 1003
Non-zero elements in train set: 0.09439083

```

Finished testing with dictionary size: 1000

Testing with dictionary size: 3000

```

[1] "<START> i saw this movie about a year ago and found it"
Train data dimensions: 25000 3003
Non-zero elements in train set: 0.03836119

```

Finished testing with dictionary size: 3000

Testing with dictionary size: 5000

```

[1] "<START> i saw this movie about a year ago and found it"
Train data dimensions: 25000 5003
Non-zero elements in train set: 0.02459676

```

Finished testing with dictionary size: 5000

Testing with dictionary size: 10000

[1] "<START> i saw this movie about a year ago and found it"

Train data dimensions: 25000 10003

Non-zero elements in train set: 0.01316756

Finished testing with dictionary size: 10000